

Load-Robust, Physics-Informed Machine Learning for Stator Inter-Turn Fault Diagnosis in Induction Motors: A Leakage-Safe Cascade for Detection, Severity, and Faulted-Phase Identification

Shashwat Garg

Abstract—Stator winding turn-to-turn (inter-turn) short-circuit faults are a leading cause of induction-motor failure, and early, load-robust diagnosis from non-invasive stator-current measurements remains an open problem. Most reported machine-learning (ML) studies achieve near-perfect in-distribution accuracy but never test generalization to unseen operating points, rarely release code, and seldom use leakage-safe evaluation. This paper presents a complete, openly released, physics-informed framework that performs hierarchical diagnosis from the three-phase stator current—fault-onset detection, healthy/faulty classification, ordinal severity estimation (percentage of shorted turns), and faulted-phase identification—and that is evaluated with a strict group-wise protocol and, above all, with *Leave-One-Load-Out* (LOLO) cross-validation. A 357-case PSCAD/EMTDC corpus ($f_s = 25$ kHz, $f_0 = 50$ Hz) is segmented into 18,033 labelled windows described by 72 physics-grounded features. LOLO reveals a clean structural result: detection (macro-F1 0.95) and faulted-phase identification (macro-F1 0.997) are intrinsically load-robust because they ride on load-invariant negative-sequence signatures, whereas absolute severity is load-dependent. We then (i) benchmark *eight* models on identical splits—including the tabular foundation model *TabPFN-2.5* and an *xLSTM* sequence network—establishing that, in-distribution, the task is saturated (all feature models within ± 0.01 macro-F1) so that the honest discriminator is cross-load behaviour; and (ii) introduce and ablate four load-robustness mechanisms: a sequence-component Stockwell tensor (cross-load severity within-one $0.75 \rightarrow 1.0$), a load-invariance-regularized feature selection, a FiLM load-conditioned multi-task network (cross-load severity $0.93 \rightarrow 0.98$), and a physics-anchored domain adaptation that recovers a simulation-trained detector on a public real-motor dataset ($0.48 \rightarrow 0.98$ macro-F1 with two labelled real samples). The deployed cascade reaches within-one-level accuracy of 0.95 in-distribution and 0.92 across unseen loads, with calibrated detection (expected calibration error 0.021), graceful noise degradation, and 100% fault-onset detection at 0% pre-fault false alarm. A within-hardware evaluation of all eight models on the real motor confirms detection transfers (macro-F1 up to 0.99) and severity is usable; per-window faulted-phase localization is the main gap (macro-F1 0.72 vs. 0.97 in simulation) but per-recording aggregation recovers it to 0.93–0.95—all reported transparently. Finally, leakage-safe Optuna tuning with a cross-load objective and physics-consistent, signature-preserving training-fold augmentation, evaluated under 10-fold grouped cross-validation, further raise cross-load severity (within-one $0.91 \rightarrow 0.94$; in-distribution $0.93 \rightarrow 0.98$; mean absolute error -20%) and detection (LOLO $0.96 \rightarrow 0.98$). All

code, splits, and artifacts are released.

Index Terms—Induction motor, inter-turn fault, stator winding, fault diagnosis, severity estimation, machine learning, negative-sequence current, leakage-safe evaluation, foundation model, transfer learning.

I. INTRODUCTION

A. Motivation

INDUCTION motors (IMs) drive the majority of industrial processes, and stator winding insulation degradation is among their most common and consequential failure modes. An incipient inter-turn (turn-to-turn) short circuit (ITSC) shorts a small fraction of the turns of one phase winding; the circulating fault current accelerates local heating and, if undetected, escalates to a phase-to-ground or phase-to-phase fault and catastrophic machine loss [1], [2]. Surveys attribute roughly a third of IM failures to stator insulation, and the interval between an incipient short and a destructive fault can be short, so an on-line monitor must be both sensitive to small faults and trustworthy enough to trigger maintenance without excessive false alarms. Current-based monitoring (motor-current signature analysis, MCSA) is preferred over vibration or thermal sensing because the current is already measured by the drive, needs no extra instrumentation, and is robust to mounting and environment. Three practical questions must be answered—*is* a fault present, *how severe* is it, and *which phase* is affected—and, critically, the answers must hold up as the mechanical load changes.

B. Literature Review

Penman *et al.* [1] established the inter-turn sideband frequencies in the stator current, and Cardoso *et al.* [2] introduced the Extended Park's Vector Approach (EPVA), in which the $2f_0$ component of the current-vector modulus scales with winding asymmetry. The negative-sequence current and negative-sequence impedance are widely used asymmetry descriptors [4], [5], and phasor compensation suppresses the baseline asymmetry due to supply unbalance. Mazzoletti *et al.* [6] use the positive-sequence fifth-harmonic current during soft-starter start-up to distinguish ITSC from high-resistance connections; Singh and Shaik [7] combine the Stockwell transform with support-vector machines to detect, type, and localize stator faults. Data-driven studies feed

Manuscript received May 26, 2026. (Corresponding author: Shashwat Garg.)

The author is with the Department of Electrical Engineering (affiliation placeholder), e-mail: author@institution.edu.

Code, dataset splits, and result artifacts are publicly available at <https://github.com/Regal-s/induction-motor-fault-diagnosis>.

short-time-Fourier amplitudes of symmetrical components to SVM/MLP classifiers [8], compare ANN/LSTM/physics-informed networks on a public dataset [9] (observing that intermediate severities overlap), or ensemble convolutional networks on spectrograms [10]. Recent architectures—temporal-convolutional [31], attention/Transformer [30], selective state-space [20], and Kolmogorov–Arnold networks [21], as well as tabular foundation models—are largely demonstrated on bearing vibration or generic tabular data rather than three-phase stator current. A comprehensive review [3] flags three systemic weaknesses of the field: suspiciously perfect accuracies from leaky train/test partitions, almost no released code or standardized datasets, and little validation across operating conditions. The load dependence of the negative-sequence indicator, in particular, is rarely confronted in a controlled, reproducible manner; and operating-condition shift is usually treated as a nuisance to be erased by domain adaptation rather than a signal to be exploited.

C. Summary of Contributions

This paper treats ITSC diagnosis as a hierarchical problem (onset $\rightarrow$ detection $\rightarrow$ severity $\rightarrow$ faulted phase) and adopts *Leave-One-Load-Out* (LOLO) evaluation as its central methodological device. Concretely, the delivered system and its contributions are:

- A *leakage-safe evaluation framework*—grouping by operating point and using LOLO as the primary metric—together with a SHAP/Fisher analysis that explains *why* detection and faulted phase are load-invariant while absolute severity is load-dependent. We argue this should be standard practice for the field.
- A *four-stage cascade* on 72 physics-grounded features: onset detection, healthy/faulty detection, ordinal severity estimation, and faulted-phase identification, with calibrated probabilities, conformal severity intervals, and an online onset trigger.
- A *controlled eight-model benchmark* on identical splits—logistic regression, SVM-RBF, k -NN, MLP, random forest, XGBoost, the tabular foundation model *TabPFN*-2.5, and an *xLSTM* sequence network—showing the in-distribution task is saturated and that cross-load behaviour, not in-distribution accuracy, is the meaningful axis of comparison.
- Four *load-robustness mechanisms*, each ablated under LOLO: a sequence-component Stockwell tensor (SCST) representation, a load-invariance-regularized feature selection (LIR-mRMR), a phase-coupled FiLM *load-conditioned* multi-task network (PCM-Net) that injects rather than erases the operating point, and physics-anchored domain adaptation (PADA) validated simulation-to-experiment, plus an honest negative result on physics-constrained generative augmentation.
- A *within-hardware evaluation* of all eight models on a real motor (leakage-safe, by recording and by repetition), establishing that detection transfers (macro-F1 up to 0.99) and severity is usable, and transparently identifying faulted-phase localization as the main simulation-to-reality gap.




Fig. 1. Phase-A current (top) and per-cycle RMS envelope (bottom) of a faulted record, with motor start (2.0 s), load change (4.0 s), fault inception (7.2 s), and clearing (8.0 s). The fault interval supplies faulty windows; the pre-fault loaded region supplies matched-load healthy windows.

- A fully *reproducible* release of code, leakage-safe splits, feature tables, and per-model result artifacts.

A unifying theme emerges: the mechanical load is the crux of stator-fault diagnosis, and every proposed mechanism independently improves *cross-load* performance.

II. PROPOSED METHODOLOGY

A. Overview

The system ingests a short window of three-phase current and produces, in sequence, an estimate of the fault-onset instant, a healthy/faulty decision, and—when faulty—an ordinal severity level and the faulted phase. Fig. 1 shows the underlying signal: each faulted record follows a fixed schedule (motor start, load change, fault inception, clearing) verified from the per-cycle current envelope. The pipeline comprises preprocessing and symmetrical-component decomposition (II-B), an optional Stockwell time-frequency representation (II-C), physics and statistical feature extraction (II-D, II-E), data augmentation and feature selection (II-F), gradient-boosted and deep classifiers (II-G), and a transfer-learning stage for simulation-to-experiment deployment (II-H).

The hierarchical, gated organization is deliberate: the four questions (is a fault present, of what severity, in which phase, and when did it start) are nested, have different decision boundaries, and have different load sensitivities, so a single flat multi-class model would conflate them and waste the structure. Gating the severity and phase stages on the detector means those stages are trained and operated only on data they are meant to explain, and lets per-stage errors be attributed and audited independently. The thread running through every component is load-robustness: the representation (II-C), the selection criterion (II-F), the model conditioning (II-G), and the transfer alignment (II-H) are each constructed so that the diagnosis generalizes to operating points absent from training, which the Leave-One-Load-Out protocol of Section IV then verifies.

B. Signal Preprocessing and Symmetrical-Component Decomposition

Each recording is segmented into windows of $W = 2000$ samples (four 50 Hz cycles); a 75% overlap is used inside

the short 0.8 s fault interval and a sparser stride elsewhere. Crucially, the *pre-fault* loaded region of every faulted record (genuinely healthy, since the fault is inserted only at 7.2 s) is labelled healthy, providing matched-load healthy data at the same loads as the faulty windows and preventing the detector from separating classes by load level. To remove load-induced amplitude while preserving the inter-phase imbalance that encodes the fault, all three phases of a window are divided by a single per-recording scale; magnitude features are additionally divided by the within-window positive-sequence fundamental $|I_1|$ to obtain load-invariant variants. The Clarke transform underlying the EPVA and the Park ellipse is

$$i_\alpha = \frac{2}{3} \left(i_a - \frac{1}{2}i_b - \frac{1}{2}i_c \right), \quad i_\beta = \frac{1}{\sqrt{3}} (i_b - i_c), \quad (1)$$

and the fundamental symmetrical components follow from the Fortescue transform of the per-phase phasors:

$$\begin{bmatrix} \mathbf{I}_0 \\ \mathbf{I}_1 \\ \mathbf{I}_2 \end{bmatrix} = \frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 1 & a & a^2 \\ 1 & a^2 & a \end{bmatrix} \begin{bmatrix} \mathbf{I}_a \\ \mathbf{I}_b \\ \mathbf{I}_c \end{bmatrix}, \quad a = e^{j2\pi/3}. \quad (2)$$

The negative-sequence ratio $|I_2|/|I_1|$ grows monotonically with the shorted-turn fraction and is the primary severity descriptor, while the angle $\angle I_2$ separates into three $\sim 120^\circ$ -spaced clusters that identify the faulted phase.

C. Stockwell-Based Time-Frequency Analysis

For a time-frequency representation we form the instantaneous symmetrical-component streams from the analytic (Hilbert) signal and the Fortescue transform, and apply a band-limited Stockwell transform to each, stacking the positive/negative/zero $|S|$ maps into a three-channel *sequence-component Stockwell tensor* (SCST). The negative-sequence channel carries the load-invariant inter-turn signature, and the frequency-dependent resolution of the Stockwell transform tracks the $2f_0$ EPVA line and slip sidebands better than a fixed-window spectrogram. Because symmetrical-component *magnitudes* are phase-agnostic, faulted-phase localization uses the negative-sequence *angle* rather than the magnitude tensor (Section IV-E).

D. Time-Domain and Spectral Characteristics

Physically grounded descriptors capture the fault signature. The EPVA severity factor is

$$\text{SF} = \frac{A(2f_0)}{\text{DC}} \text{ of } |\mathbf{i}_P(t)| = \sqrt{i_\alpha^2 + i_\beta^2}, \quad (3)$$

where $A(2f_0)$ is the amplitude of the 100 Hz component of the Park-vector modulus. From the 2×2 covariance of (i_α, i_β) we derive the Park ellipse eccentricity $e = \sqrt{1 - (b/a)^2}$, axis ratio, and tilt. Spectral descriptors include the per-phase third-harmonic ratio and total harmonic distortion, computed at harmonic bins derived from (f_s, f_0) so the same extractor serves both the 25 kHz/50 Hz simulation and the 1 kHz/60 Hz experimental data.

TABLE I
FEATURE GROUPS (72 FEATURES PER WINDOW).

Group	Features
Symmetrical comp.	$ I_0 , I_1 , I_2 $; $ I_2 / I_1 $, $ I_0 / I_1 $; $\angle I_1, \angle I_2$ (+ sin/cos); compensated I_2^c
EPVA / Park	severity factor; Park-modulus mean/std; ellipse axes, eccentricity, axis ratio, tilt
Harmonics	third-harmonic ratio and THD (per phase, mean)
Per-phase stats	RMS, peak, crest, std, skew, kurtosis ($\times 3$) and inter-phase differences
Load-invariant	$ I_1 $ -normalized variants of all magnitude features

E. Statistical Feature Extraction

Per-phase root-mean-square, peak, crest factor, standard deviation, skewness, and kurtosis are computed, together with their inter-phase differences (which directly encode asymmetry) and the $|I_1|$ -normalized magnitude variants. With the symmetrical-component, EPVA, ellipse, and harmonic quantities this yields the 72-feature bank summarized in Table I. Severity is encoded as an ordinal rank $r \in \{0, \dots, 6\}$ over $\{0.3, 0.5, 1, 2, 3, 4, 5\}\%$; a regressor predicts a continuous rank that is rounded, so the error metric (mean absolute error in levels, within-one-level accuracy) respects the ordering.

F. Data Augmentation and Feature Selection

Feature selection. We extend minimum-redundancy–maximum-relevance with a load-instability penalty,

$$J(f) = \text{Rel}(f) - \lambda \text{Red}(f) - \gamma \text{LoadInstability}(f), \quad (4)$$

where $\text{LoadInstability}(f)$ is the across-load drift of the feature’s class-conditional mean normalized by its global spread ($\gamma=0$ recovers plain mRMR). LIR-mRMR selects dimensionless, load-stable features and, on the load-dependent severity task, generalizes better at compact size (Section IV-E). *Augmentation.* For the deployed pipeline we apply, to the *training fold only*, a set of physics-consistent augmentations selected from a survey of time-series and signature-guided methods [25], [26]: joint amplitude (load) scaling, per-channel SNR jitter, mild magnitude-warp, phase-coherent time shift, and—for the faulted-phase task—exact faulted-phase rotation (a cyclic channel permutation with the negative-sequence angle rotated $\mp 120^\circ$). Every amplitude operation acts *jointly* on the three phases so the inter-phase imbalance (the fault signature) and Kirchhoff balance $i_a + i_b + i_c = 0$ are preserved; a current-balance reject filter guards this, and signature-corrupting operations (per-channel scaling, segment permutation, phase-spectrum randomization) are excluded. Augmented windows are re-featurised through the same extractor. We separately study a conditional generator of three-phase current with a hard Kirchhoff projection and a soft negative-sequence-versus-severity loss; the hard constraint is enforced exactly, with an honest negative result on controllability reported in Section IV-E. Hyperparameters of the gradient-boosted models are selected by Optuna TPE [24] with a *cross-load* inner objective (leave-one-load-out over the training loads, the

model-selection principle of domain generalization [23]), with early stopping on an inner held-out load.

G. Models: Gradient Boosting, Foundation Model, and Deep Networks

Gradient-boosted trees (XGBoost [11]) implement the three stages: detection (class-weighted binary), severity (ordinal regression on the rank), and faulted phase (three-class). Stages 1 and 3 use the full feature set; the severity stage uses the load-invariant subset and the *measured* load. The detector gates the downstream stages: windows predicted faulty are passed to the severity and phase models. To place the cascade in context we evaluate, on identical splits, six classical feature classifiers and two recent learners on the same 72 features or raw windows: *TabPFN*-2.5 [12], a pretrained tabular foundation model that performs in-context Bayesian inference over a feature table, and *xLSTM* [13], an extended LSTM with exponential gating and a stabilizer state applied to the raw three-phase current. As a load-conditioned deep alternative we develop *PCM-Net*, a phase-coupled, multi-task network with a dilated-temporal backbone (a compact stand-in for a selective state-space model) and *feature-wise linear modulation* (FiLM) [27] driven by the measured operating point, so one model spans loads instead of erasing the condition; a current-balance physics term regularizes training. A fault-onset stage thresholds a per-cycle negative-sequence index $\rho(k) = |I_2(k)|/|I_1(k)|$ against a trailing pre-fault reference, $\rho(k) > \bar{\rho} + 6\sigma$ for three consecutive cycles. Detection probabilities are assessed with expected calibration error and the Brier score [17], and severity uncertainty with conformalized quantile regression [15], [16].

H. Transfer Learning

To deploy a simulation-trained model on a real machine we use physics-anchored domain adaptation (PADA). A naively transferred model collapses under domain shift; we therefore align feature distributions across domains, anchoring the alignment on the domain-stable physics features (negative-sequence ratio, angles, $|I_1|$ -normalized magnitudes), and optionally calibrate with a few labelled target repetitions. Section IV-E quantifies the recovery on a public real-motor dataset.

III. EXPERIMENTAL SETUP AND DATASET PREPARATION

A. Description of the Experimental Setup

The primary dataset is generated from a three-phase squirrel-cage induction motor with stator inter-turn faults modelled in PSCAD/EMTDC ($f_0 = 50$ Hz, $f_s = 25$ kHz, 10 s per run, 500 samples per cycle). The fault is introduced through a tapped stator winding that short-circuits a controlled fraction μ of one phase's turns, and each run follows a fixed schedule that mimics a realistic duty cycle: the motor is energized at 2.0 s, the mechanical load is applied at 4.0 s, the inter-turn short is inserted at 7.2 s, and it is cleared ≈ 0.8 s later (Fig. 1). This schedule lets a single run furnish both a faulted segment and a matched-load healthy segment, and it exercises the start-up inrush, the loaded steady state, and the fault transient. The three-phase currents are exported and used

directly (no filtering), so the diagnosis operates on the raw drive measurement.

For transfer-learning validation we additionally use a public real-motor dataset [28]: a 0.75 hp squirrel-cage IM, three-phase current acquired with split-core current transformers at 1 kHz ($f_0 = 60$ Hz, no mechanical load), with a healthy class and inter-turn shorts at 10/20/30/40% of turns in each of the three phases over five repetitions per class. The two sources differ in fundamental frequency, sampling rate, severity scale, sensing chain, and the presence of a healthy baseline, providing a demanding cross-domain benchmark; both are mapped to a common window/label schema and the feature extractor derives its harmonic bins from (f_s, f_0) so that one pipeline serves both.

B. Data Sets Generated from the Experimental Setup

The simulation corpus spans severity $\mu \in \{0.3, 0.5, 1, 2, 3, 4, 5\}\%$, faulted phase $\in \{A, B, C\}$, load $\in \{\text{NL}, 20, 40, 60, 80, 100\}\%$, and several fault-inception instants, with a healthy set over the full load range: 357 cases (336 faulty, 21 healthy). Segmentation yields 18,033 labelled windows (11,760 faulty, 6,273 healthy), balanced across severity and phase. The real-motor set yields 1,235 windows. Windows are grouped by operating point (severity $\times$ phase $\times$ load), with all inception variants of a point kept together, so that no window from a record can appear in both training and test.

IV. PERFORMANCE EVALUATION

A. Implementation and Evaluation Protocol

Features are computed in Python (NumPy/Pandas) and models trained with scikit-learn and XGBoost (400–500 trees, depth 5, learning rate 0.05, subsample/colsample 0.8); deep models use PyTorch and *TabPFN*-2.5 runs on CPU. Two evaluation protocols are used: Stratified Group k -Fold ($k=5$, grouped by operating point) for the in-distribution estimate, and *Leave-One-Load-Out* (LOLO)—all windows at one load form the test set—as the primary metric, measuring generalization to an operating point never seen in training. Reported metrics are accuracy, macro-F1, severity within-one-level accuracy, mean absolute error (MAE) in levels, quadratic-weighted kappa (QWK), and per-phase F1, computed on out-of-fold predictions and reported as a mean over folds. Stages 2–3 are trained on faulty windows but evaluated on the windows the detector predicts faulty, so cascade numbers reflect realistic error propagation; predictions are also aggregated per recording (median severity, majority-vote phase). Scalars are fit on the training partition only, every benchmark model sees identical folds, no per-model tuning is done on test folds, and all randomized components use fixed seeds.

B. Headline Results

Table II summarizes the deployed cascade. *The central finding is structural rather than numerical*: in-distribution, every stage exceeds 0.97, but only LOLO separates the three tasks into those that genuinely generalize across operating

TABLE II
DEPLOYED CASCADE: PER-STAGE AND END-TO-END PERFORMANCE (MEAN OVER FOLDS).

Stage	Metric	Group k -Fold	LOLO (primary)
1 Detection	macro-F1 / ROC-AUC	0.977 / 0.994	0.948
2 Severity (load-aware)	within-1 / QWK	0.997 / 0.987	0.946
3 Phase ID	macro-F1	0.972	0.997
End-to-end	within-1 / per-record	0.949 / 0.927	0.916 / 0.918

TABLE III
SEVERITY WITHIN-ONE-LEVEL ACCURACY UNDER FEATURE CHOICE. THE LOAD-AWARE MODEL IS THE DEPLOYED CONFIGURATION.

Severity feature set	Group k -Fold	LOLO	No-load
All features (magnitudes incl.)	0.993	0.887	0.39
Load-invariant + $ I_1 $ -norm	0.990	0.729	0.80
+ measured load (proposed)	0.997	0.946	0.95

points (detection, macro-F1 0.948; faulted phase, macro-F1 0.997) and the one that does not without intervention (absolute severity). The detection ROC has area 0.994 (Fig. 2) and the faulted-phase confusion matrix (Fig. 3) is near-diagonal with phase C perfect. End-to-end, scoring a window correct when healthy windows are detected healthy and faulty windows have the correct phase and a severity within one level, the gated cascade reaches within-one accuracy 0.949 in-distribution and 0.916 across unseen loads (per-recording 0.927 and 0.918); Fig. 4 contrasts the two protocols. Notably, faulted-phase identification is the only metric for which cross-load *exceeds* in-distribution, a direct consequence of its load-invariance.

C. Why the Tasks Differ: Cross-Load Analysis and the Load-Aware Fix

SHAP attribution shows detection and phase ride on the *negative-sequence* ratio and angle and on inter-phase RMS differences—all load-invariant ratios—whereas an initial severity model built on absolute magnitudes *collapsed* at the no-load extreme under LOLO (within-one 0.39), because the $|I_2|$ magnitude that encodes severity scales with the operating point. Two complementary remedies recover it (Table III). First, restricting severity to dimensionless, load-invariant features and $|I_1|$ -normalized magnitudes lifts the no-load case from 0.39 to 0.80. Second—and decisively—supplying the *measured* load (a sensor input, not a label) raises pooled LOLO within-one from 0.73 to 0.95, since most held-out loads are interpolated between neighbours. Severity errors are then almost entirely between adjacent levels (Fig. 5, QWK 0.987), which validates the ordinal treatment. The per-held-out-load breakdown (Table IV) localizes the residual difficulty to the light-load extremes (no-load and 20%), where a low-severity fault current most overlaps the healthy distribution—a difficulty visible *only* under LOLO.

D. Eight-Model Benchmark

Table V and Fig. 6 benchmark eight models on *identical* leakage-safe splits, spanning classical feature classifiers, a tabular foundation model, and a sequence deep network. The

TABLE IV
PER-HELD-OUT-LOAD DETECTION (LOLO, MACRO-F1). DIFFICULTY CONCENTRATES AT LIGHT LOADS.

	Hold-out load	NL	20%	40%	60%	80%	100%
Detection F1	0.885	0.749	1.000	1.000	1.000	1.000	1.000

first observation is that the *in-distribution task is saturated*: the six feature classifiers fall within ± 0.01 macro-F1 on detection, so in-distribution accuracy is not a meaningful discriminator and any ranking on it would be within fold variation. The honest comparison is therefore cross-load. There, three results stand out. (1) The tabular foundation model *TabPFN-2.5* is the strongest detector and uniquely load-robust, scoring macro-F1 0.979 in-distribution *and* 0.979 cross-load, whereas gradient boosting drops from 0.976 to 0.948; it also gives the sharpest in-distribution severity (0.997 within-one, MAE 0.033 level, the lowest of all models) and degrades gracefully across loads (0.969). (2) Severity is the task that separates models: tree ensembles extrapolate poorly to unseen loads (random forest LOLO within-one 0.688), while the MLP and SVM remain robust (0.982, 0.975). (3) Faulted phase is load-robust for nearly every model (LOLO macro-F1 ≥ 0.96), confirming that the negative-sequence-angle signature is model-agnostic. The raw-signal *xLSTM* trails the feature learners on every task (LOLO macro-F1 0.890 detection, within-one 0.908 severity, 0.961 phase) yet improves on the earlier raw-signal deep baselines (1-D ResNet and PCM-Net, 0.898 detection; Table VI), reaffirming the data-efficiency of physically engineered features over raw-waveform deep nets in this regime. The feature pipeline is also lightweight—feature extraction plus XGBoost inference costs ≈ 0.34 ms per window on CPU (Fig. 7), and TabPFN-2.5 needs no task-specific training—so the practitioner can choose foundation-model robustness or gradient-boosted speed without sacrificing accuracy.

Statistical significance. To test whether these differences are real rather than fold noise we follow the Demšar protocol [22]: a Friedman omnibus test across the seven feature/foundation models on the per-fold (group k -fold, five blocks) and per-load (LOLO, six blocks) primary scores, with Holm-corrected pairwise Wilcoxon signed-rank tests (Table VII). Two conclusions follow. First, *in-distribution detection and phase identification show no significant model differences* ($p = 0.26$ and $p = 0.76$), which quantitatively confirms that the in-distribution task is saturated—the choice among feature models cannot be justified on in-distribution accuracy. Second, the differences that matter *are* significant: absolute severity differs across models under both protocols ($p = 0.0013$ in-distribution, $p = 0.0008$ cross-load) and detection differs across loads ($p = 0.050$), driven by the gradient-boosted model’s cross-load drop and the tree ensembles’ poor severity extrapolation (Random Forest LOLO within-one 0.70). By average Friedman rank, TabPFN-2.5 and the MLP are strongest on cross-load severity and the tree ensembles weakest. With only five-to-six blocks, however, no *pairwise* comparison survives Holm correction (smallest adjusted $p = 0.19$); the omnibus is therefore the reliable signal, and localizing pairwise differences would require more folds




Fig. 2. Detection. Left: receiver operating characteristic, overall and per held-out load (overall AUC = 0.994). Right: confusion matrix (group k -fold, out-of-fold).

TABLE V

FULL MODEL BENCHMARK ON IDENTICAL LEAKAGE-SAFE SPLITS: STRATIFIEDGROUPKFOLD (GK, IN-DISTRIBUTION) AND LEAVE-ONE-LOAD-OUT (LOLO, CROSS-LOAD, PRIMARY). DETECTION AND PHASE REPORT MACRO-F1; SEVERITY REPORTS WITHIN-ONE-LEVEL ACCURACY. BEST PER COLUMN IN BOLD. TABPFN-2.5 AND XLSTM ARE ADDED IN THIS WORK.

Model	Detection (macro-F1)		Severity (within-1)		Phase (macro-F1)	
	GK	LOLO	GK	LOLO	GK	LOLO
Logistic Regression	0.978	0.978	0.997	0.971	0.970	0.920
SVM-RBF	0.978	0.977	0.991	0.975	0.984	0.967
k -NN	0.977	0.977	0.937	0.942	0.982	0.981
MLP	0.978	0.978	0.997	0.982	0.956	0.981
Random Forest	0.978	0.979	0.985	0.688	0.965	0.997
XGBoost	0.976	0.948	0.990	0.943	0.976	0.997
TabPFN-2.5	0.979	0.979	0.997	0.969	0.968	0.991
xLSTM	0.957	0.890	0.982	0.908	0.921	0.961




Fig. 3. Faulted-phase confusion (group k -fold, out-of-fold). Phase C is perfect; only minor A $\leftrightarrow$ B confusion remains. Cross-load (LOLO) phase accuracy is 0.997.




Fig. 4. Per-stage and end-to-end performance, in-distribution (group k -fold) vs. cross-load (LOLO). Phase identification is the only metric for which cross-load exceeds in-distribution, reflecting its load-invariance.

or additional operating points—a limitation we flag rather than over-claim from.

E. Load-Robustness Mechanisms (Ablations)

Table VIII ablates the four proposed mechanisms under LOLO, the protocol on which each is designed to help. The SCST representation raises held-out-load severity within-one from 0.75 (raw-phase Stockwell) to 1.0; consistent with the physics, its magnitude tensor cannot localize phase, which is why phase uses the negative-sequence angle instead. LIR-




Fig. 5. Severity confusion (group k -fold). Errors are almost entirely between adjacent levels, justifying the ordinal treatment and the within-one-level metric (QWK = 0.987).

TABLE VI
FEATURE-BASED CASCADE VS. A RAW-SIGNAL DEEP BASELINE.

Task / metric	XGBoost (features)	ResNet-1D (raw)
Detection F1 (group k -fold)	0.977	0.898
Detection F1 (LOLO, load 20%)	0.749	0.401
Phase macro-F1 (group k -fold)	0.972	0.848
Severity within-1 (group k -fold)	0.990	1.000

TABLE VII
SIGNIFICANCE ACROSS THE SEVEN FEATURE/FOUNDATION MODELS (FRIEDMAN OMNIBUS ON PER-FOLD / PER-LOAD PRIMARY SCORES; DETECTION AND PHASE USE MACRO-F1, SEVERITY WITHIN-ONE-LEVEL ACCURACY).

Task	Protocol	χ^2	p	Outcome
Detection	group k -fold	7.66	0.264	n.s.
Detection	LOLO	12.61	0.050	sig.
Severity	group k -fold	21.84	0.0013	sig.
Severity	LOLO	22.85	0.0008	sig.
Phase	group k -fold	3.35	0.764	n.s.
Phase	LOLO	6.00	0.423	n.s.

mRMR beats plain *mRMR* at compact size ($k=10$ within-one 0.848 vs. 0.740) with roughly half the cross-load feature drift, yielding a sparse, interpretable subset. *PCM-Net*'s FiLM load-conditioning lifts LOLO severity within-one from 0.932 to **0.982**—the best in the study, exceeding even the load-aware gradient-boosted model (0.946); FiLM does not help the load-invariant phase task, exactly as the physics predicts, which is itself evidence the mechanism does what it claims. *PADA* closes the simulation-to-experiment gap on the public real-motor dataset: a naively transferred detector scores macro-F1 0.48 (it collapses to predicting all-faulty under covariate shift), physics-anchored unsupervised alignment reaches 0.50 on the hard phase task, and few-shot calibration with two real repetitions recovers detection to **0.98**. Finally, the physics-constrained generator enforces current balance exactly (residual 0.000 vs. 0.991 for an unconstrained generator) but

TABLE VIII
LOAD-ROBUSTNESS MECHANISMS (LOLO).

Mechanism	Cross-load result
SCST representation	severity within-1 0.75 $\rightarrow$ 1.0; phase needs angle
LIR-mRMR selection	$k=10$ severity within-1 0.740 $\rightarrow$ 0.848; drift halved
PCM-Net + FiLM	severity within-1 0.932 $\rightarrow$ 0.982 (best; > load-aware GBM 0.946)
PADA (sim $\rightarrow$ real)	detection F1 0.48 $\rightarrow$ 0.98 (few-shot); 0.50 phase (unsup.)
PC-Diff	current-balance exact; generation not controllable (negative)

TABLE IX
NOISE ROBUSTNESS (TRAIN CLEAN, TEST NOISY; GROUP k -FOLD).

SNR	Detect macro-F1	Severity within-1	Phase acc.
clean	0.977	0.986	0.958
40 dB	0.927	0.980	0.954
30 dB	0.854	0.975	0.945
20 dB	0.792	0.971	0.939

TABLE X
FAULT-ONSET DETECTION (SAMPLED CASES).

Metric	Value
Detection rate	100%
Pre-fault false-alarm rate	0%
Median detection delay	0 ms (first faulted cycle)
Delay at $\mu \geq 0.5\%$ / $\mu = 0.3\%$	0 ms / ≈ 1.4 s

did not yield severity-controllable samples under our compute budget—an honest negative result that we report to motivate a stronger conditional generator rather than omit.

F. Calibration, Onset, Noise, and Training Behaviour

Stage-1 probabilities are well calibrated (ECE = 0.021, Brier = 0.021), so the detector's confidence can gate maintenance actions directly. Conformalized quantile regression improves severity-interval coverage over naive quantiles, with group-conditional calibration noted as the remedy for residual under-coverage. The onset detector (Table X) flags 100% of faults with no pre-fault false alarms; for severity $\geq 0.5\%$ the fault is caught within the first faulted cycle, and only the smallest 0.3% faults incur a lag. Under additive noise (Table IX), severity and phase are essentially flat to 20 dB SNR while detection is the noise-sensitive stage—expected, since separating a small fault from healthy operation is hardest in noise. Training is stable for both the gradient-boosted and deep models (Fig. 8), and the per-stage precision/recall/F1 breakdown is given in Fig. 9.

G. Hardware Validation (Within-Domain)

All results so far are from simulation; we now train and test the same models *on a real motor*. Using the public ibarram dataset [28] (a 0.75 hp IM, three-phase current at 1 kHz, no load; 65 recordings = 13 classes $\times$ 5 repetitions), we evaluate every model within-hardware under two leakage-safe protocols: StratifiedGroupKFold grouped by recording,




Fig. 6. Model benchmark across the three diagnostic tasks under both protocols (GroupKFold vs. LOLO). Hatched bars denote the two models added in this work (TabPFN-2.5, xLSTM). Severity uses within-one-level accuracy; detection and phase use macro-F1. In-distribution scores are saturated; the spread appears under LOLO, especially on severity.




Fig. 7. Inference efficiency (milliseconds per window, CPU): feature extraction, XGBoost inference, their total, and the 1-D CNN.

and Leave-One-Repetition-Out (LORO), which holds out an entire unseen recording per class. Severity uses the dataset’s own four-level scale (10/20/30/40%). Table XI reports the outcome and Fig. 10 the XGBoost confusion matrices, and we deliberately report the weaknesses as fully as the strengths.

Three findings stand out. First, *detection transfers to hardware*: TabPFN-2.5 and the gradient-boosted/forest models reach macro-F1 0.95–0.99 on the real motor, and the LORO score (an entirely unseen recording) is as strong as the in-domain one, so fault *presence* is decided reliably from real current. Second, *severity is moderate*: within-one-level accuracy is 0.83–0.96 but exact accuracy falls to 0.6–0.75 (TabPFN-2.5 best, with the lowest mean absolute error of 0.34 level), reflecting genuinely harder discrimination of adjacent shorted-turn fractions on real signals. Third, *per-window faulted-phase localization is weak on hardware* (macro-F1 0.64–0.72, versus 0.97 in simulation), because the per-window negative-sequence *angle* estimate is noisy at the low 1 kHz rate. However, the realistic unit of diagnosis is a recording, not a single window: *aggregating a recording’s windows by majority vote recovers phase macro-F1 to 0.85 (group k -fold) and 0.93–*

TABLE XI
WITHIN-HARDWARE RESULTS ON THE REAL IBARRAM MOTOR
(STRATIFIED GROUPKFOLD BY RECORDING, GK; AND
LEAVE-ONE-REPETITION-OUT, LORO). DETECTION AND PHASE REPORT
MACRO-F1; SEVERITY REPORTS WITHIN-ONE-LEVEL ACCURACY ON ITS
FOUR-LEVEL SCALE. BEST PER COLUMN IN BOLD. PHASE VALUES ARE
PER-WINDOW; PER-RECORDING MAJORITY VOTE RAISES XGBOOST PHASE
MACRO-F1 TO 0.85 (GK) AND 0.95 (LORO).

Model	Detection		Severity (w-1)		Phase	
	GK	LORO	GK	LORO	GK	LORO
Logistic Regression	0.817	0.864	0.830	0.818	0.664	0.678
SVM-RBF	0.877	0.923	0.866	0.818	0.661	0.689
k -NN	0.829	0.884	0.827	0.804	0.642	0.648
MLP	0.846	0.880	0.866	0.811	0.664	0.690
Random Forest	0.923	0.947	0.900	0.829	0.650	0.718
XGBoost	0.966	0.955	0.900	0.836	0.636	0.707
TabPFN-2.5	0.994	0.954	0.929	0.864	0.643	0.701
xLSTM	0.708	0.699	0.963	0.957	0.651	0.648

0.95 (*LORO*) for the gradient-boosted and forest models—a +0.22 gain under LORO that requires no new features. A controlled ablation found that reference-invariant feature engineering (relative angle, per-phase asymmetry) did *not* improve beyond this, which localizes the cause to per-window estimator variance rather than feature design; closing the residual per-window gap (higher sampling, supervised phase calibration on the target machine) is left to future work. Consistent with the simulation benchmark, TabPFN-2.5 is again the strongest tabular model, and the raw-signal xLSTM trails on detection and phase (its high severity within-one reflects adjacent-level predictions rather than exact accuracy).

H. Hyperparameter Tuning and Data Augmentation

We finally isolate the effect of leakage-safe hyperparameter tuning and physics-consistent training-fold augmentation on the deployed XGBoost pipeline, evaluated under 10-fold StratifiedGroupKFold (the 147 operating-point groups satisfy $n_{\text{folds}} \leq n_{\text{groups}}$) and LOLO. Hyperparameters are tuned for cross-load generalization (Section II-F) and frozen for both protocols; augmentation is applied to the training fold only and




Fig. 8. Training behaviour. Left: XGBoost detection convergence (train/validation log-loss vs. boosting round). Right: 1-D CNN learning curves (training loss and validation macro-F1 vs. epoch), showing stable convergence without over-fitting.




Fig. 9. Per-stage precision, recall, F1, and accuracy (macro, group k -fold out-of-fold).

TABLE XII

EFFECT OF LEAKAGE-SAFE TUNING AND PHYSICS-CONSISTENT TRAINING-FOLD AUGMENTATION ON THE DEPLOYED XGBOOST PIPELINE (10-FOLD STRATIFIEDGROUPKFOLD AND LOLO). DETECTION/PHASE: MACRO-F1; SEVERITY: WITHIN-ONE-LEVEL ACCURACY. BEST PER ROW IN BOLD.

Task (metric)	Protocol	Default	Tuned	Tuned+Aug
Detection (macro-F1)	10-fold GK	0.976	0.976	0.979
Detection (macro-F1)	LOLO	0.964	0.964	0.977
Severity (within-1)	10-fold GK	0.934	0.937	0.976
Severity (within-1)	LOLO	0.911	0.908	0.935
Phase (macro-F1)	10-fold GK	0.964	0.936	0.964
Phase (macro-F1)	LOLO	0.994	0.997	0.997

re-featurised. Table XII reports the three conditions. Tuning alone is marginal on the already-saturated metrics (consistent with the significance analysis of Section IV-D), but *augmentation is the decisive lever for cross-load generalization*: it raises severity within-one-level accuracy from 0.911 to 0.935 under LOLO and from 0.934 to 0.976 in-distribution, cutting the mean absolute error by 20% (0.917 $\rightarrow$ 0.731 level), and lifts detection LOLO macro-F1 from 0.964 to 0.977, while faulted-phase identification remains at 0.997. This confirms the augmentation principle: synthesising training windows toward unseen operating points—by jointly scaling the three phases, so the inter-phase imbalance and current balance are preserved—directly improves cross-load diagnosis without fabricating the fault signature.

I. Discussion: Positioning and Deployment

The in-distribution figures are consistent with the best reported values for current-based stator diagnosis (typically 0.96–1.00 accuracy [3], [8], [9]); the contribution here is not a higher in-distribution number—which the saturated benchmark shows is no longer informative—but a demonstration of *which* of those numbers survive a change of operating point, and a set of mechanisms that make the load-dependent ones survive. The in-distribution-versus-LOLO severity gap (within-one 0.99 vs. 0.39 before the load-aware fix) quantifies how much a conventional evaluation would have over-promised; we therefore advocate cross-condition (LOLO) reporting, alongside interpretability evidence, as standard practice. From a deployment standpoint the pipeline is attractive: it consumes only the already-measured three-phase current, runs in well under a millisecond per window on a CPU, exposes interpretable physical quantities (the negative-sequence ratio and EPVA factor that drive a decision), provides calibrated probabilities with an online onset trigger suitable for a protection relay, and offers a foundation-model option (TabPFN-2.5) that is both the most load-robust detector and training-free. Within-hardware validation on a real motor (Section IV-G) confirms that detection transfers (macro-F1 up to 0.99) and severity is usable (within-one ≥ 0.86); per-window faulted-phase localization is weak (macro-F1 0.72) but per-recording aggregation—the realistic diagnostic unit—recovers it to 0.93–0.95. The remaining limitations are that the load-robustness results are validated only in simulation (the public real dataset has no load variation), and that a full instrumented campaign with supply-voltage unbalance and richer real severities remains future work.

V. CONCLUSION

We presented a physics-informed, leakage-safe machine-learning framework for stator inter-turn fault diagnosis in induction motors that performs onset detection, healthy/faulty detection, ordinal severity estimation, and faulted-phase identification from the three-phase stator current. By adopting Leave-One-Load-Out evaluation we showed that detection




Fig. 10. Within-hardware confusion matrices on the real motor (XGBoost, GroupKFold out-of-fold): detection (left), four-level severity (middle), faulted phase (right). Detection is clean and severity errors are near-diagonal, but the phase matrix shows substantial off-diagonal confusion—the main simulation-to-reality gap.

and faulted-phase identification are intrinsically load-robust, while absolute severity is load-dependent and is recovered to within-one accuracy of 0.95 across unseen loads once the model is made load-aware. A controlled eight-model benchmark—including the tabular foundation model TabPFN-2.5 and an xLSTM sequence network—demonstrated that the in-distribution task is saturated and that cross-load behaviour is the meaningful axis, on which TabPFN-2.5 is the most load-robust detector. Four mechanisms (a sequence-component Stockwell tensor, a load-invariance-regularized feature selection, a FiLM-conditioned multi-task network, and physics-anchored domain adaptation) each independently improved cross-load performance, and the framework was validated simulation-to-experiment on a public real-motor dataset. The end-to-end cascade attains within-one accuracy of 0.95 in-distribution and 0.92 across unseen loads, with calibrated detection, graceful noise degradation, and reliable onset detection. A within-hardware evaluation of all eight models on the real motor confirmed that detection transfers (macro-F1 up to 0.99) and severity is usable; per-window faulted-phase localization is the main gap (macro-F1 0.72) but per-recording aggregation recovers it to 0.93–0.95. Future work includes closing the residual per-window phase gap (higher sampling, supervised phase calibration), a stronger conditional generator for physics-constrained augmentation (the signature-preserving training-fold augmentation introduced here already cut cross-load severity error by 20%), full selective-state-space and Kolmogorov–Arnold backbones, an experimental campaign with supply voltage-unbalance rejection, and group-conditional conformal calibration of the severity intervals. All code, dataset splits, and result artifacts are released for reproducibility.

REFERENCES

- [1] J. Penman, H. G. Sedding, B. A. Lloyd, and W. T. Fink, “Detection of inter-turn short circuits in the stator windings of operating motors,” *IEEE Trans. Energy Convers.*, vol. 9, no. 4, pp. 652–658, 1994.
- [2] A. J. M. Cardoso, S. M. A. Cruz, and D. S. B. Fonseca, “Inter-turn stator winding fault diagnosis in three-phase induction motors, by Park’s vector approach,” *IEEE Trans. Energy Convers.*, vol. 14, no. 3, pp. 595–598, 1999.
- [3] R. Hussain, M. Alshaikh Saleh, and S. S. Refaat, “Various faults classification of industrial application of induction motors using supervised machine learning: A comprehensive review,” *IEEE Access*, vol. 13, pp. 146649–146675, 2025.
- [4] D. García-Pérez *et al.*, “Machine learning for inverter-fed motors monitoring and fault detection: An overview,” *IEEE Access*, vol. 12, 2024.
- [5] S. Ruzimov *et al.*, “Detection of inter-turn short-circuit faults for inverter-fed induction motors based on negative-sequence current analysis,” *Sensors*, vol. 25, no. 15, art. 4844, 2025.
- [6] M. A. Mazzeletti, G. R. Bossio, P. D. Donolo, and M. Donolo, “Turn-to-turn short-circuit fault and high-resistance connection diagnosis for induction motor during start-up transient,” *IEEE Trans. Ind. Appl.*, vol. 61, no. 2, pp. 1927–1935, 2025.
- [7] M. Singh and A. G. Shaik, “Incipient fault detection in stator windings of an induction motor using Stockwell transform and SVM,” *IEEE Trans. Instrum. Meas.*, vol. 69, no. 12, pp. 9496–9504, 2020.
- [8] P. Pietrzak and M. Wolkiewicz, “Machine learning-based stator current data-driven PMSM stator winding fault diagnosis,” *Sensors*, vol. 22, no. 24, art. 9668, 2022.
- [9] T. Aasar *et al.*, “Predicting stator winding short-circuit faults in induction motors using machine learning: A comparative study,” *Discover Computing*, vol. 29, art. 103, 2026.
- [10] U. Ali, U. Ramzan, W. Ali, and K. A. Al-Jaafari, “An improved fault diagnosis strategy for induction motors using weighted probability ensemble deep learning,” *IEEE Access*, vol. 13, pp. 106958–106973, 2025.
- [11] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proc. ACM SIGKDD*, 2016, pp. 785–794.
- [12] N. Hollmann *et al.*, “Accurate predictions on small data with a tabular foundation model,” *Nature*, vol. 637, pp. 319–326, 2025.
- [13] M. Beck *et al.*, “xLSTM: Extended long short-term memory,” in *Proc. NeurIPS*, 2024.
- [14] S. M. Lundberg and S.-I. Lee, “A unified approach to interpreting model predictions,” in *Proc. NeurIPS*, 2017, pp. 4765–4774.
- [15] Y. Romano, E. Patterson, and E. Candès, “Conformalized quantile regression,” in *Proc. NeurIPS*, 2019.
- [16] A. N. Angelopoulos and S. Bates, “A gentle introduction to conformal prediction and distribution-free uncertainty quantification,” *arXiv:2107.07511*, 2021.
- [17] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, “On calibration of modern neural networks,” in *Proc. ICML*, 2017, pp. 1321–1330.
- [18] T. Ince, S. Kiranyaz, L. Eren, M. Askar, and M. Gabbouj, “Real-time motor fault detection by 1-D convolutional neural networks,” *IEEE Trans. Ind. Electron.*, vol. 63, no. 11, pp. 7067–7075, 2016.
- [19] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proc. IEEE CVPR*, 2016, pp. 770–778.
- [20] A. Gu and T. Dao, “Mamba: Linear-time sequence modeling with selective state spaces,” *arXiv:2312.00752*, 2023.
- [21] Z. Liu *et al.*, “KAN: Kolmogorov–Arnold networks,” in *Proc. ICLR*, 2025.
- [22] J. Demšar, “Statistical comparisons of classifiers over multiple data sets,” *J. Mach. Learn. Res.*, vol. 7, pp. 1–30, 2006.

- [23] I. Gulrajani and D. Lopez-Paz, “In search of lost domain generalization,” in *Proc. ICLR*, 2021.
- [24] T. Akiba, S. Sano, T. Yanase, T. Ohta, and M. Koyama, “Optuna: A next-generation hyperparameter optimization framework,” in *Proc. ACM SIGKDD*, 2019, pp. 2623–2631.
- [25] B. K. Iwana and S. Uchida, “An empirical survey of data augmentation for time series classification with neural networks,” *PLOS ONE*, vol. 16, no. 7, e0254841, 2021.
- [26] H. Zhang, M. Cissé, Y. N. Dauphin, and D. Lopez-Paz, “mixup: Beyond empirical risk minimization,” in *Proc. ICLR*, 2018.
- [27] E. Perez, F. Strub, H. de Vries, V. Dumoulin, and A. Courville, “FiLM: Visual reasoning with a general conditioning layer,” in *Proc. AAAI*, 2018.
- [28] M. Ibarra-Manzano *et al.*, “ITSC: Inter-turn short-circuit fault dataset for three-phase induction motors,” GitHub repository, <https://github.com/ibarram/ITSC>.
- [29] R. Cunha, “MIT short-circuit flux and current signals dataset,” Kaggle, 2021, <https://www.kaggle.com/datasets/rebecacunha/mit-short-circuit-flux-and-current-signals>.
- [30] A. Vaswani *et al.*, “Attention is all you need,” in *Proc. NeurIPS*, 2017.
- [31] S. Bai, J. Z. Kolter, and V. Koltun, “An empirical evaluation of generic convolutional and recurrent networks for sequence modeling,” *arXiv:1803.01271*, 2018.
- [32] A. Sharma *et al.*, “Recurrence-plot-based severity assessment of stator inter-turn faults using ensemble learning,” *Measurement*, 2023.