

Load-Robust, Physics-Informed Machine Learning for Stator Inter-Turn Fault Diagnosis in Induction Motors: A Leakage-Safe Cascade for Detection, Severity, and Faulted-Phase Identification

Shashwat Garg

Abstract—Stator winding turn-to-turn (inter-turn) short-circuit faults are a leading cause of induction-motor failure, and early, load-robust diagnosis from non-invasive stator-current measurements remains an open problem. Most reported machine-learning (ML) studies achieve near-perfect in-distribution accuracy but never test generalization to unseen operating points, rarely release code, and seldom use leakage-safe evaluation. This paper presents a complete, openly released, physics-informed framework that performs hierarchical diagnosis from the three-phase stator current—fault-onset detection, healthy/faulty classification, ordinal severity estimation (percentage of shorted turns), and faulted-phase identification—and that is evaluated with a strict group-wise protocol and, above all, with *Leave-One-Load-Out* (LOLO) cross-validation. A 357-case PSCAD/EMTDC corpus ($f_s = 25$ kHz, $f_0 = 50$ Hz) is segmented into 18,033 labelled windows described by 72 physics-grounded features. LOLO reveals a clean structural result: detection (macro-F1 0.95) and faulted-phase identification (macro-F1 0.997) are intrinsically load-robust because they ride on load-invariant negative-sequence signatures, whereas absolute severity is load-dependent. We then (i) benchmark *eight* models on identical splits—including the tabular foundation model *TabPFN-2.5* and an *xLSTM* sequence network—establishing that, in-distribution, the task is saturated (all feature models within ± 0.01 macro-F1) so that the honest discriminator is cross-load behaviour; and (ii) introduce and ablate four load-robustness mechanisms: a sequence-component Stockwell tensor (cross-load severity within-one $0.75 \rightarrow 1.0$), a load-invariance-regularized feature selection, a FiLM load-conditioned multi-task network (cross-load severity $0.93 \rightarrow 0.98$), and a physics-anchored domain adaptation that recovers a simulation-trained detector on a public real-motor dataset ($0.48 \rightarrow 0.98$ macro-F1 with two labelled real samples). The deployed cascade reaches within-one-level accuracy of 0.95 in-distribution and 0.92 across unseen loads, with calibrated detection (expected calibration error 0.021), graceful noise degradation, and 100% fault-onset detection at 0% pre-fault false alarm. A within-hardware evaluation of all eight models on the real motor confirms detection transfers (macro-F1 up to 0.99) and severity is usable; per-window faulted-phase localization is the main gap (macro-F1 0.72 vs. 0.97 in simulation) but per-recording aggregation recovers it to 0.93–0.95—all reported transparently. Finally, leakage-safe Optuna tuning with a cross-load objective and physics-consistent, signature-preserving training-fold augmentation, evaluated under 10-fold grouped cross-validation, further raise cross-load severity (within-one $0.91 \rightarrow 0.94$; in-distribution $0.93 \rightarrow 0.98$; mean absolute error -20%) and detection (LOLO $0.96 \rightarrow 0.98$). A novel

block of biomedical anomaly features (signal-complexity, entropy, fractal-dimension, and heart-rate-variability-style cycle-to-cycle measures imported from ECG/EEG analysis) likewise lifts cross-load detection ($+0.03$) and severity ($+0.04$) where absolute physics magnitudes do not transfer. All code, splits, and artifacts are released.

Index Terms—Induction motor, inter-turn fault, stator winding, fault diagnosis, severity estimation, machine learning, negative-sequence current, leakage-safe evaluation, foundation model, transfer learning.

I. INTRODUCTION

A. Motivation

INDUCTION motors (IMs) drive the majority of industrial processes, and stator winding insulation degradation is among their most common and consequential failure modes. An incipient inter-turn (turn-to-turn) short circuit (ITSC) shorts a small fraction of the turns of one phase winding; the circulating fault current accelerates local heating and, if undetected, escalates to a phase-to-ground or phase-to-phase fault and catastrophic machine loss [1], [2]. Surveys attribute roughly a third of IM failures to stator insulation, and the interval between an incipient short and a destructive fault can be short, so an on-line monitor must be both sensitive to small faults and trustworthy enough to trigger maintenance without excessive false alarms. Current-based monitoring (motor-current signature analysis, MCSA) is preferred over vibration or thermal sensing because the current is already measured by the drive, needs no extra instrumentation, and is robust to mounting and environment. Three practical questions must be answered—*is* a fault present, *how severe* is it, and *which phase* is affected—and, critically, the answers must hold up as the mechanical load changes.

B. Literature Review

Penman *et al.* [1] established the inter-turn sideband frequencies in the stator current, and Cardoso *et al.* [2] introduced the Extended Park's Vector Approach (EPVA), in which the $2f_0$ component of the current-vector modulus scales with winding asymmetry. The negative-sequence current and negative-sequence impedance are widely used asymmetry descriptors [4], [5], and phasor compensation suppresses the baseline asymmetry due to supply unbalance. Mazzeletti

Manuscript received May 26, 2026. (Corresponding author: Shashwat Garg.)

The author is with the Department of Electrical Engineering (affiliation placeholder), e-mail: author@institution.edu.

Code, dataset splits, and result artifacts are publicly available at <https://github.com/Regal-s/induction-motor-fault-diagnosis>.

et al. [6] use the positive-sequence fifth-harmonic current during soft-starter start-up to distinguish ITSC from high-resistance connections; Singh and Shaik [7] combine the Stockwell transform with support-vector machines to detect, type, and localize stator faults. Data-driven studies feed short-time-Fourier amplitudes of symmetrical components to SVM/MLP classifiers [8], compare ANN/LSTM/physics-informed networks on a public dataset [9] (observing that intermediate severities overlap), or ensemble convolutional networks on spectrograms [10]. Recent architectures—temporal-convolutional [38], attention/Transformer [37], selective state-space [20], and Kolmogorov–Arnold networks [21], as well as tabular foundation models—are largely demonstrated on bearing vibration or generic tabular data rather than three-phase stator current. A comprehensive review [3] flags three systemic weaknesses of the field: suspiciously perfect accuracies from leaky train/test partitions, almost no released code or standardized datasets, and little validation across operating conditions. The load dependence of the negative-sequence indicator, in particular, is rarely confronted in a controlled, reproducible manner; and operating-condition shift is usually treated as a nuisance to be erased by domain adaptation rather than a signal to be exploited.

C. Summary of Contributions

This paper treats ITSC diagnosis as a hierarchical problem (onset $\rightarrow$ detection $\rightarrow$ severity $\rightarrow$ faulted phase) and adopts *Leave-One-Load-Out* (LOLO) evaluation as its central methodological device. Concretely, the delivered system and its contributions are:

- A *leakage-safe evaluation framework*—grouping by operating point and using LOLO as the primary metric—together with a SHAP/Fisher analysis that explains *why* detection and faulted phase are load-invariant while absolute severity is load-dependent. We argue this should be standard practice for the field.
- A *four-stage cascade* on 72 physics-grounded features: onset detection, healthy/faulty detection, ordinal severity estimation, and faulted-phase identification, with calibrated probabilities, conformal severity intervals, and an online onset trigger.
- A *controlled eight-model benchmark* on identical splits—logistic regression, SVM-RBF, k -NN, MLP, random forest, XGBoost, the tabular foundation model *TabPFN*-2.5, and an *xLSTM* sequence network—showing the in-distribution task is saturated and that cross-load behaviour, not in-distribution accuracy, is the meaningful axis of comparison.
- A *novel biomedical anomaly-feature* block—signal-complexity, entropy, fractal-dimension, and heart-rate-variability-style cycle-to-cycle measures imported from ECG/EEG/HRV analysis, computed on the current residual and Park-vector modulus—that complements the physics features and improves cross-load detection (+0.03) and severity (+0.04) where absolute-magnitude features do not transfer.
- Four *load-robustness mechanisms*, each ablated under LOLO: a sequence-component Stockwell tensor

(SCST) representation, a load-invariance-regularized feature selection (LIR-mRMR), a phase-coupled FiLM *load-conditioned* multi-task network (PCM-Net) that injects rather than erases the operating point, and physics-anchored domain adaptation (PADA) validated simulation-to-experiment, plus an honest negative result on physics-constrained generative augmentation.

- A *within-hardware evaluation* of all eight models on a real motor (leakage-safe, by recording and by repetition), establishing that detection transfers (macro-F1 up to 0.99) and severity is usable, and transparently identifying faulted-phase localization as the main simulation-to-reality gap.
- A fully *reproducible* release of code, leakage-safe splits, feature tables, and per-model result artifacts.

A unifying theme emerges: the mechanical load is the crux of stator-fault diagnosis, and every proposed mechanism independently improves *cross-load* performance.

II. PROPOSED METHODOLOGY

A. Overview

The system ingests a short window of three-phase current and produces, in sequence, an estimate of the fault-onset instant, a healthy/faulty decision, and—when faulty—an ordinal severity level and the faulted phase. Fig. 2 shows the underlying signal: each faulted record follows a fixed schedule (motor start, load change, fault inception, clearing) verified from the per-cycle current envelope. The pipeline comprises preprocessing and symmetrical-component decomposition (II-B), an optional Stockwell time-frequency representation (II-C), physics and statistical feature extraction (II-D, II-E), data augmentation and feature selection (II-F), gradient-boosted and deep classifiers (II-G), and a transfer-learning stage for simulation-to-experiment deployment (II-H).

The hierarchical, gated organization is deliberate: the four questions (is a fault present, of what severity, in which phase, and when did it start) are nested, have different decision boundaries, and have different load sensitivities, so a single flat multi-class model would conflate them and waste the structure. Gating the severity and phase stages on the detector means those stages are trained and operated only on data they are meant to explain, and lets per-stage errors be attributed and audited independently. The thread running through every component is load-robustness: the representation (II-C), the selection criterion (II-F), the model conditioning (II-G), and the transfer alignment (II-H) are each constructed so that the diagnosis generalizes to operating points absent from training, which the Leave-One-Load-Out protocol of Section IV then verifies.

Step-by-step diagnosis pipeline. The end-to-end inference path, executed once per arriving 80 ms window of three-phase current $\mathbf{i}(t) = [i_a, i_b, i_c]^T$, is summarized below; each step also defines the corresponding training-time module. Fig. 1 renders the same flow graphically.

- **S1. Acquisition and indexing.** Read the three phase currents from the simulation export (or the live ADC); every recording is indexed against its operating-point label

(severity, phase, load, inception) and its verified region bounds (motor start, load change, fault inception, clearing) using the per-cycle RMS envelope (Fig. 2).

- S2. *Windowing and labelling.* Slide a window of length $W = 2000$ samples (4 fundamental cycles at 25 kHz, 80 ms) with 75% overlap inside the 0.8 s fault interval and a sparser stride elsewhere. The *pre-fault* loaded region of every faulted record is labelled healthy, yielding matched-load healthy data at every load.
- S3. *Per-recording normalization.* Divide all three phases of the window by a *single* scale equal to the RMS of the case's healthy region across the three phases; this removes load-induced amplitude while preserving inter-phase imbalance, the actual fault signature.
- S4. *Symmetrical-component decomposition.* Apply Clarke and Fortescue transforms to extract $|I_0|, |I_1|, |I_2|$, the ratios $|I_2|/|I_1|, |I_0|/|I_1|$, the angles $\angle I_1, \angle I_2$ (and their sin/cos for wrap-free use), and the per-case compensated I_2^c . The ratio $|I_2|/|I_1|$ grows monotonically with shorted-turn fraction; $\angle I_2$ clusters $\sim 120^\circ$ apart by faulted phase.
- S5. *Physics, spectral and statistical features.* Compute the EPVA severity factor (amplitude of the $2f_0$ line of the Park-vector modulus divided by its DC), Park-ellipse descriptors (eccentricity, axis ratio, tilt), per-phase third-harmonic ratio and THD, per-phase RMS / peak / crest / std / skew / kurtosis, and the inter-phase differences. Magnitude features are also divided by the within-window $|I_1|$ to obtain the n_* *load-invariant* variants.
- S6. *Biomedical anomaly features.* On the residual (current minus its fitted fundamental) and on the Park-vector modulus, compute HRV cycle-to-cycle waveform distance and Poincaré SD1/SD2, Hjorth mobility and complexity, permutation / spectral / sample / dispersion entropy, and Higuchi / Katz / Petrosian fractal dimension—scale-free measures that quantify how regular the waveform is rather than which spectral line moves.
- S7. *Leakage-safe splitting and training-fold augmentation.* Windows are grouped by operating point (severity $\times$ phase $\times$ load, with all inception variants kept together) so that no window from one recording can appear in both train and test; physics-consistent augmentations (joint amplitude scaling, per-channel SNR jitter, magnitude warp, phase-coherent shift, faulted-phase rotation) are applied to the training fold only.
- S8. *Stage 0—fault-onset trigger.* A per-cycle test $\rho(k) = |I_2(k)|/|I_1(k)| > \bar{\rho} + 6\sigma$ for three consecutive cycles raises the onset flag; this is essentially free to compute and gates the cascade in deployment.
- S9. *Stage 1—detection (healthy vs. faulty).* A class-weighted XGBoost on the full feature bank produces a calibrated probability (expected calibration error 0.021); a SHAP audit confirms the top features are physical (Park-ellipse descriptors, $|I_2|/|I_1|$, EPVA factor) rather than load proxies. If *healthy*, the cascade terminates; if *faulty*, S10 and S11 execute in parallel.
- S10. *Stage 2—severity (ordinal).* The load-invariant feature subset is augmented with the *measured* load and an ordinal regressor predicts the severity rank $r \in \{0, \dots, 6\}$

Stator Inter-Turn Fault Diagnosis Pipeline (Track A, deployed)




Fig. 1. Block diagram of the deployed diagnosis pipeline (steps S1–S12 of §II-A). Three-phase stator current enters at the top; per-recording normalization (S3) and the 96-feature extractor (S4–S6) feed Stage 0 (onset trigger on the per-cycle $|I_2|/|I_1|$ index) and the gated XGBoost cascade—Stage 1 detection → {Stage 2 severity (load-aware ordinal), Stage 3 faulted phase (3-class)}. Per-recording aggregation (S12) produces the final verdict.

over $\{0.3, 0.5, 1, 2, 3, 4, 5\}\%$. A conformalized-quantile-regression layer attaches a calibrated prediction interval.

- S11. *Stage 3—faulted phase (A/B/C).* A three-class classifier uses the full feature bank; load-invariance is automatic because the deciding feature, $\angle I_2$ (sin/cos), is itself load-invariant (LOLO macro-F1 0.997).
- S12. *Per-recording aggregation.* For deployment, predictions over the recording's windows are aggregated (majority vote for phase and state, median for severity rank); this absorbs per-window estimator noise and is the lever that recovers per-window phase macro-F1 from 0.72 to 0.93–0.95 on real hardware.

B. Signal Preprocessing and Symmetrical-Component Decomposition

Each recording is segmented into windows of $W = 2000$ samples (four 50 Hz cycles); a 75% overlap is used inside the short 0.8 s fault interval and a sparser stride elsewhere. Crucially, the *pre-fault* loaded region of every faulted record (genuinely healthy, since the fault is inserted only at 7.2 s) is labelled healthy, providing matched-load healthy data at the same loads as the faulty windows and preventing the detector from separating classes by load level. To remove load-induced




Fig. 2. Phase-A current (top) and per-cycle RMS envelope (bottom) of a faulted record, with motor start (2.0 s), load change (4.0 s), fault inception (7.2 s), and clearing (8.0 s). The fault interval supplies faulty windows; the pre-fault loaded region supplies matched-load healthy windows.

amplitude while preserving the inter-phase imbalance that encodes the fault, all three phases of a window are divided by a single per-recording scale; magnitude features are additionally divided by the within-window positive-sequence fundamental $|I_1|$ to obtain load-invariant variants. The Clarke transform underlying the EPVA and the Park ellipse is

$$i_\alpha = \frac{2}{3} \left(i_a - \frac{1}{2} i_b - \frac{1}{2} i_c \right), \quad i_\beta = \frac{1}{\sqrt{3}} (i_b - i_c), \quad (1)$$

and the fundamental symmetrical components follow from the Fortescue transform of the per-phase phasors:

$$\begin{bmatrix} \mathbf{I}_0 \\ \mathbf{I}_1 \\ \mathbf{I}_2 \end{bmatrix} = \frac{1}{3} \begin{bmatrix} 1 & 1 & 1 \\ 1 & a & a^2 \\ 1 & a^2 & a \end{bmatrix} \begin{bmatrix} \mathbf{I}_a \\ \mathbf{I}_b \\ \mathbf{I}_c \end{bmatrix}, \quad a = e^{j2\pi/3}. \quad (2)$$

The negative-sequence ratio $|I_2|/|I_1|$ grows monotonically with the shorted-turn fraction and is the primary severity descriptor, while the angle $\angle I_2$ separates into three $\sim 120^\circ$ -spaced clusters that identify the faulted phase.

Window length: why four cycles, not one. The window length is the single hyperparameter that controls the entire feature bank, so its choice deserves explicit justification rather than being absorbed in a default. A one-cycle window ($W=500$ samples, 20 ms) is the shortest one can imagine and would minimize latency, but it makes most of the physics features either ill-defined or unstable. A W -sample window has spectral resolution

$$\Delta f = \frac{f_s}{W} = \frac{1}{T_{\text{win}}}, \quad (3)$$

so $W=500$ at $f_s=25$ kHz gives $\Delta f = 50$ Hz, the spacing between the fundamental and the $2f_0$ EPVA line. The DC, the 50 Hz fundamental, the 100 Hz EPVA line, and the early Penman sidebands $f_0[1 \pm n(1-s)/p]$ all collapse into a few overlapping bins, and the EPVA severity factor of (3)—defined precisely as $A(2f_0)/\text{DC}$ of $|\mathbf{i}_P(t)|$ and the strongest single severity scalar in the bank—cannot be measured. With four cycles ($W=2000$, $T_{\text{win}}=80$ ms) the resolution improves to $\Delta f = 12.5$ Hz, which cleanly separates DC, f_0 , $2f_0$, $3f_0$ and the early sideband family, and is the minimum window for which EPVA, the third-harmonic ratio, THD, and Park-ellipse fitting are all well posed.

TABLE I
WINDOW-LENGTH TRADE-OFF AT $f_s = 25$ kHz. Δf IS THE SPECTRAL RESOLUTION OF (3). N_{FAULT} IS THE TYPICAL NUMBER OF NON-OVERLAPPING WINDOWS PER FAULTED RECORDING IN THE 0.8 s FAULT INTERVAL. “EPVA / HARMONICS USABLE” REQUIRES $\Delta f \ll 2f_0$.

Cycles	W	T_{win}	Δf	N_{fault}	EPVA / harm.
1	500	20 ms	50.0 Hz	40	no (bins merge)
4	2000	80 ms	12.5 Hz	10	yes (chosen)
10	5000	200 ms	5.0 Hz	4	yes, but redundant
20	10,000	400 ms	2.5 Hz	2	yes; too few examples

Four cycles is also the shortest window for which the higher-order statistical features (skewness, kurtosis, Park-modulus mean / std) and the biomedical block (cycle-to-cycle waveform distance, Poincaré SD1/SD2, fractal dimension) are meaningful: skew and kurtosis are notoriously high-variance on small samples, and the HRV-style descriptors are by definition cross-cycle and need at least four heartbeats to form a Poincaré plot. Numerically, the trade-off against longer windows is sample efficiency: the fault interval is only ≈ 0.8 s (20,000 samples) long, so $W=2000$ with stride 500 produces ~ 36 windows per faulted recording and a balanced training set of 11,760 faulted plus 6,273 healthy windows, whereas $W=5000$ (10 cycles) yields only ~ 8 non-overlapping windows per recording and a more redundant corpus, and $W=10,000$ (20 cycles) gives too few examples to support cross-validation. Table I makes the trade-off explicit. The combination of (i) sufficient spectral resolution to expose the $2f_0$ EPVA signature, (ii) stable higher-order statistics, (iii) dense sampling of the short fault interval, and (iv) deployment latency below one line cycle on a Jetson-class CPU (< 1 ms feature extraction at 80 ms acquisition latency) all point to $W = 2000$, which is why it was retained over the broader $\{1, 4, 10, 20\}$ -cycle sweep envisaged in the original plan.

C. Stockwell-Based Time-Frequency Analysis

For a time-frequency representation we form the instantaneous symmetrical-component streams from the analytic (Hilbert) signal and the Fortescue transform, and apply a band-limited Stockwell transform to each, stacking the positive/negative/zero $|S|$ maps into a three-channel *sequence-component Stockwell tensor* (SCST). The negative-sequence channel carries the load-invariant inter-turn signature, and the frequency-dependent resolution of the Stockwell transform tracks the $2f_0$ EPVA line and slip sidebands better than a fixed-window spectrogram. Because symmetrical-component *magnitudes* are phase-agnostic, faulted-phase localization uses the negative-sequence *angle* rather than the magnitude tensor (Section IV-E).

D. Time-Domain and Spectral Characteristics

Physically grounded descriptors capture the fault signature. The EPVA severity factor is

$$\text{SF} = \frac{A(2f_0)}{\text{DC}} \text{ of } |\mathbf{i}_P(t)| = \sqrt{i_\alpha^2 + i_\beta^2}, \quad (4)$$

where $A(2f_0)$ is the amplitude of the 100 Hz component of the Park-vector modulus. From the 2×2 covariance of (i_α, i_β)

TABLE II
FEATURE GROUPS (72 FEATURES PER WINDOW).

Group	Features
Symmetrical comp.	$ I_0 , I_1 , I_2 $; $ I_2 / I_1 $, $ I_0 / I_1 $; $\angle I_1, \angle I_2$ (+ sin/cos); compensated I_2^c
EPVA / Park	severity factor; Park-modulus mean/std; ellipse axes, eccentricity, axis ratio, tilt
Harmonics	third-harmonic ratio and THD (per phase, mean)
Per-phase stats	RMS, peak, crest, std, skew, kurtosis ($\times 3$) and inter-phase differences
Load-invariant	$ I_1 $ -normalized variants of all magnitude features
Biomedical (novel)	HRV cycle-to-cycle waveform distance & Poincaré SD1/SD2; Hjorth mobility/complexity; permutation/spectral/sample/dispersion entropy; Higuchi/Katz/Petrosian fractal dimension (on the residual & Park modulus)

we derive the Park ellipse eccentricity $e = \sqrt{1 - (b/a)^2}$, axis ratio, and tilt. Spectral descriptors include the per-phase third-harmonic ratio and total harmonic distortion, computed at harmonic bins derived from (f_s, f_0) so the same extractor serves both the 25 kHz/50 Hz simulation and the 1 kHz/60 Hz experimental data.

E. Statistical Feature Extraction

Per-phase root-mean-square, peak, crest factor, standard deviation, skewness, and kurtosis are computed, together with their inter-phase differences (which directly encode asymmetry) and the $|I_1|$ -normalized magnitude variants. With the symmetrical-component, EPVA, ellipse, and harmonic quantities this yields the 72-feature bank summarized in Table II. Severity is encoded as an ordinal rank $r \in \{0, \dots, 6\}$ over $\{0.3, 0.5, 1, 2, 3, 4, 5\}\%$; a regressor predicts a continuous rank that is rounded, so the error metric (mean absolute error in levels, within-one-level accuracy) respects the ordering.

F. Biomedical Anomaly Features

Beyond the physics descriptors we introduce a block of *biomedical anomaly-detection features* transferred from ECG/EEG/heart-rate-variability (HRV) analysis—a view absent from the motor-fault literature, which treats the current as a deterministic sum of sinusoids and asks *which spectral line moves*, rather than *how regular* the waveform is. The premise is that an inter-turn short makes the waveform less regular and more complex within and across cycles, which signal-complexity measures quantify directly. Crucially, these are computed on the *residual* (current minus its fitted fundamental) and the Park-vector modulus—not the raw current, whose clean fundamental is trivially predictable. Three groups are used: (i) an *HRV-on-current* block that treats each fundamental cycle as a heartbeat—cycle-to-cycle waveform distance to the median cycle and Poincaré descriptors (SD1, SD2, SD1/SD2) of the per-cycle RMS series, the direct analog of HRV [32], [33] and the most physically faithful import, since the fault *is* cycle-to-cycle irregularity; (ii) complexity scalars—Hjorth mobility and complexity [29], permutation [28] and

spectral entropy, sample [27] and dispersion entropy [31]; and (iii) fractal dimension—Higuchi [30], Katz and Petrosian. The complexity and fractal members are amplitude-scale-free, which (Table XIV) aids cross-load generalization. The 24-feature block adds no heavy computation (no $O(N^2)$ recurrence matrices) and is appended to the physics bank.

G. Data Augmentation and Feature Selection

Feature selection. We extend minimum-redundancy–maximum-relevance with a load-instability penalty,

$$J(f) = \text{Rel}(f) - \lambda \text{Red}(f) - \gamma \text{LoadInstability}(f), \quad (5)$$

where $\text{LoadInstability}(f)$ is the across-load drift of the feature’s class-conditional mean normalized by its global spread ($\gamma=0$ recovers plain mRMR). LIR-mRMR selects dimensionless, load-stable features and, on the load-dependent severity task, generalizes better at compact size (Section IV-E). *Augmentation.* For the deployed pipeline we apply, to the *training fold only*, a set of physics-consistent augmentations selected from a survey of time-series and signature-guided methods [25], [26]: joint amplitude (load) scaling, per-channel SNR jitter, mild magnitude-warp, phase-coherent time shift, and—for the faulted-phase task—exact faulted-phase rotation (a cyclic channel permutation with the negative-sequence angle rotated $\mp 120^\circ$). Every amplitude operation acts *jointly* on the three phases so the inter-phase imbalance (the fault signature) and Kirchhoff balance $i_a + i_b + i_c = 0$ are preserved; a current-balance reject filter guards this, and signature-corrupting operations (per-channel scaling, segment permutation, phase-spectrum randomization) are excluded. Augmented windows are re-featurised through the same extractor. We separately study a conditional generator of three-phase current with a hard Kirchhoff projection and a soft negative-sequence-versus-severity loss; the hard constraint is enforced exactly, with an honest negative result on controllability reported in Section IV-E. Hyperparameters of the gradient-boosted models are selected by Optuna TPE [24] with a *cross-load* inner objective (leave-one-load-out over the training loads, the model-selection principle of domain generalization [23]), with early stopping on an inner held-out load.

H. Models: Gradient Boosting, Foundation Model, and Deep Networks

Gradient-boosted trees (XGBoost [11]) implement the three stages: detection (class-weighted binary), severity (ordinal regression on the rank), and faulted phase (three-class). Stages 1 and 3 use the full feature set; the severity stage uses the load-invariant subset and the *measured* load. The detector gates the downstream stages: windows predicted faulty are passed to the severity and phase models. To place the cascade in context we evaluate, on identical splits, six classical feature classifiers and two recent learners on the same 72 features or raw windows: *TabPFN-2.5* [12], a pretrained tabular foundation model that performs in-context Bayesian inference over a feature table, and *xLSTM* [13], an extended LSTM with exponential gating and a stabilizer state applied to the raw three-phase current. As a load-conditioned deep alternative we develop *PCM-Net*,

a phase-coupled, multi-task network with a dilated-temporal backbone (a compact stand-in for a selective state-space model) and *feature-wise linear modulation* (FiLM) [34] driven by the measured operating point, so one model spans loads instead of erasing the condition; a current-balance physics term regularizes training. A fault-onset stage thresholds a per-cycle negative-sequence index $\rho(k) = |I_2(k)|/|I_1(k)|$ against a trailing pre-fault reference, $\rho(k) > \bar{\rho} + 6\sigma$ for three consecutive cycles. Detection probabilities are assessed with expected calibration error and the Brier score [17], and severity uncertainty with conformalized quantile regression [15], [16]. Fig. 3 contrasts the two tracks side by side: Track A (deployed) takes the 96 physics/medical features into the gated XG-Boost cascade with a per-recording aggregator, while Track B (alternative) feeds the raw three-phase window through the PCM-Net stem, phase-coupling attention, FiLM-conditioned backbone, and per-task heads.

I. Transfer Learning

To deploy a simulation-trained model on a real machine we use physics-anchored domain adaptation (PADA). A naively transferred model collapses under domain shift; we therefore align feature distributions across domains, anchoring the alignment on the domain-stable physics features (negative-sequence ratio, angles, $|I_1|$ -normalized magnitudes), and optionally calibrate with a few labelled target repetitions. Section IV-E quantifies the recovery on a public real-motor dataset.

III. EXPERIMENTAL SETUP AND DATASET PREPARATION

Two complementary datasets are used in this work. A large *simulation* corpus from PSCAD/EMTDC supports the in-distribution, leave-one-load-out, and ablation studies in §IV-B–H; a public *real-motor* dataset [35] supports within-hardware validation (§IV-J), simulation-to-experiment transfer learning (§IV-E), and the real-time edge deployment of §IV-K. The two sources are deliberately heterogeneous in machine, line frequency, sampling rate, severity scale, sensing chain, and the presence of a healthy baseline; this section documents each, the construction of the test-case matrix, and the window segmentation and labelling protocol that feeds the models.

A. Simulation Dataset: PSCAD/EMTDC Stator Inter-Turn Faults

Machine and fault model. A three-phase squirrel-cage induction motor is modelled in PSCAD/EMTDC at line frequency $f_0 = 50$ Hz. The stator inter-turn fault is implemented as a *tapped winding*: a controlled fraction μ of one phase’s turns is tapped out and short-circuited through a fault impedance, so that the model exposes both magnitude and phase asymmetry while remaining electromagnetically consistent with the unfaulted stator. The three line currents (i_a, i_b, i_c) are exported as raw time-series; no filtering or spectral shaping is applied before the feature pipeline so that the diagnosis operates on the raw measurement.

Sampling and run length. Each run is sampled at $f_s = 25$ kHz, giving 500 samples per fundamental cycle and a

TABLE III
PSCAD/EMTDC SIMULATION CORPUS—ENUMERATION OF THE
OPERATING-POINT TEST MATRIX AND RECORDING COUNTS.

Axis	Levels
Fault type	Stator inter-turn (tapped winding)
Severity μ (% shorted turns)	{0.3, 0.5, 1, 2, 3, 4, 5}
Faulted phase	{A, B, C}
Load (% rated)	{NL, 20, 40, 60, 80, 100}
Fault inception t_f (s)	{7.200, 7.202, ..., 7.220} (11 variants)
Faulted recordings ($7 \times 3 \times 6 \times 11$ – unused cells)	336
Healthy recordings	21
Total simulation recordings	357
Sampling f_s	25 kHz (500 samples/cycle)
Run length (faulted / healthy)	10 s / 5 s
Operating-point groups (CV grouping)	147

40 μ s time step. Faulted runs are 10 s long (250,000 samples); healthy runs are 5 s long (125,000 samples). Three multi-channel .out files per case carry the full measurement bus; the diagnostically relevant channels are the three line currents in kA, used directly throughout this paper.

Duty-cycle schedule. Each faulted run follows a fixed event schedule that mimics a realistic motor duty cycle (Fig. 2): the motor is energized at 2.0 s (start-up inrush settles by ~ 2.7 s), the mechanical load is applied at 4.0 s (skipped for the no-load variant), the inter-turn short is inserted at 7.2 s, and the fault is cleared ≈ 0.8 s later. This single-run schedule furnishes both a faulted segment (samples 180,000–200,000) and a *matched-load* healthy segment (the pre-fault region, samples $\sim 110,000$ –178,000), the latter providing load-matched healthy data and a per-recording baseline for negative-sequence compensation (§II-B).

Operating-point matrix. The simulation corpus enumerates a 4-axis test-case matrix:

- Severity $\mu \in \{0.3, 0.5, 1, 2, 3, 4, 5\}$ % shorted turns (7 levels).
- Faulted phase $\in \{A, B, C\}$.
- Mechanical load $\in \{\text{NL}, 20, 40, 60, 80, 100\}$ % (6 levels, NL = no load).
- Fault-inception time $t_f \in \{7.20, 7.202, 7.204, \dots, 7.220\}$ s (a baseline inception at 7.20 s plus 10 “Extra TT” variants per (μ , phase, load) triple sweeping the inception over one half-cycle of f_0 to randomise the fault-initiation angle).

A separate *healthy* subset is recorded over the full load range. The resulting counts—336 faulted recordings (126 base + 210 inception variants = $7 \times 3 \times 6 \times 11$) plus 21 healthy recordings—are summarised in Table III. Total: 357 recordings yielding 18,033 labelled windows after segmentation (Table V).

B. Real-Motor Experimental Dataset

Machine and acquisition. The real-motor validation uses a publicly released laboratory dataset [35] acquired on a 0.75 hp squirrel-cage induction motor at line frequency $f_0 = 60$ Hz. Three-phase line currents are sampled at $f_s = 1$ kHz with split-core current transformers (no mechanical load applied; the motor runs at near-rated no-load speed). The motor’s stator

Two Diagnosis Tracks: Physics-Feature Cascade (Track A, deployed) vs. PCM-Net (Track B, deep)




Fig. 3. The two diagnosis tracks evaluated in this work. **Track A** (left, the *deployed* cascade): NumPy/SciPy feature extractor $\rightarrow$ Stage 0 onset $\rightarrow$ Stage 1 XGBoost detector $\rightarrow$ {Stage 2 ordinal-severity head, Stage 3 3-class phase head} $\rightarrow$ per-recording aggregation. **Track B** (right, the deep alternative): raw $(3, W)$ window $\rightarrow$ multi-scale 1-D conv stem $\rightarrow$ phase-coupling graph attention $\rightarrow$ dilated-temporal (Mamba-surrogate) backbone $\rightarrow$ FiLM load conditioning $\rightarrow$ {detect, severity (KAN), phase} heads with a current-balance PINN regularizer. Both tracks accept the same input and emit the same verdict; Section IV-D reports their compared performance.

has a healthy and faulted winding configuration controlled by external tappings, so the same machine produces both classes under the same sensing chain—an important property for a clean within-hardware comparison.

Fault matrix. The real-motor dataset enumerates a 3-axis matrix (Table IV): *state* $\in \{\text{healthy, faulted}\}$; within faulted, *severity* $\in \{10, 20, 30, 40\}\%$ shorted turns (4 coarser levels than the simulation’s 7); *faulted phase* $\in \{A, B, C\}$; and *repetition* $\in \{01, 02, 03, 04, 05\}$ (5 replicates per condition). The resulting 13 distinct conditions (1 healthy + $4 \times 3 = 12$ faulted) $\times$ 5 repetitions yield 65 recordings, sliced into 1,235 labelled windows. Repetition 05 is held out for testing throughout this paper; train splits use 01–04.

Cross-dataset role. The two datasets are complementary rather than redundant: PSCAD/EMTDC carries the load axis (essential for LOLO and the load-aware-severity story of §IV-C) but is synthetic; the real-motor dataset carries the sensing and machine realism but no load axis. Together they support an in-distribution evaluation, a cross-load generalization test, a within-hardware test, a simulation-to-experiment domain-adaptation test, and the real-time edge-deployment test of §IV-K.

C. Window Segmentation, Labelling, and Sample Distribution

Window length is fixed at $W = 2000$ samples for the simulation corpus (4 fundamental cycles at 25 kHz, 80 ms)

TABLE IV
REAL-MOTOR DATASET [35]—ENUMERATION OF THE TEST MATRIX AND RECORDING COUNTS.

Axis	Levels
Motor	0.75 hp squirrel-cage IM, $f_0 = 60$ Hz
Sampling f_s	1 kHz
Sensing chain	Split-core current transformers
Mechanical load	None (no-load operation)
Severity (% shorted turns)	$\{10, 20, 30, 40\}$
Faulted phase	$\{A, B, C\}$
Healthy condition	$\{\text{HLT}\}$
Repetitions per condition	$\{01, 02, 03, 04, \mathbf{05}\}$
Total distinct conditions	13 (1 healthy + 12 faulted)
Total recordings	65
Labelled windows	1,235
Held-out test repetition	rep 05 (13 recordings)

and $W = 500$ samples for the real-motor dataset (~ 30 cycles at 1 kHz given the lower f_s ; 500 ms). This is the smallest window for which the EPVA $2f_0$ spectral line, the third-harmonic ratio, the Park-vector ellipse fit, and the higher-order per-phase statistics are all well-posed (the spectral resolution $\Delta f = f_s/W = 12.5$ Hz at 25 kHz cleanly separates DC, f_0 , $2f_0$, and the early Penman sidebands).

Stride and overlap. Inside the short 0.8 s simulation fault interval the stride is 500 samples (75% overlap), which is needed

TABLE V

SAMPLE DISTRIBUTION ACROSS BOTH DATASETS AFTER WINDOWING. WINDOWS ARE GROUPED BY OPERATING POINT (THE LEAKAGE-SAFE GROUPING UNIT, 147 GROUPS FOR THE SIMULATION CORPUS AND 13 FOR THE REAL-MOTOR CORPUS); INCEPTION-TIME VARIANTS AND PER-CONDITION REPETITIONS ARE BOUND INTO THE SAME GROUP AS THE BASE CASE.

	Simulation (PSCAD/EMTDC)	Real motor (5)
Recordings	357	500 (125)
– healthy	21	5 (HLT)
– faulted	336	60 (12 cond × 5 traps)
Labelled windows	18,033	1,235
– healthy	6,273	135
– faulted	11,760	1,100
Window length W	2,000 samples (80 ms)	500 samples (500 ms)
Stride (in / out of fault)	500 / 4,000	50 / 400
Severity balance	1,680 per μ level	~ 285 per μ level
Phase balance	3,920 per phase	380 per phase
CV grouping unit	operating-point (severity × phase × load)	condition × rep
Number of CV groups	147	13
Held-out test partition	per protocol (GroupKFold or LOLO)	repetition 0.5

to extract enough labelled faulted windows per recording; outside the fault interval (and across the real-motor recordings) a sparser stride is used to balance the corpus without inflating it.

Matched-load healthy labelling. The pre-fault region of every *faulted* run is labelled as a healthy window with the corresponding load attribute. This single design choice (i) furnishes load-matched healthy data at every load level, preventing the detector from learning the load distribution as a proxy for the class label, and (ii) supplies an on-recording healthy baseline used by the per-recording normalization of §II-B and by the compensated negative-sequence feature I_2^c .

Per-recording normalization. Each window is divided by a single per-recording amplitude scale (RMS over the case’s healthy region across all three phases). This removes load-induced amplitude variation while preserving the inter-phase imbalance that encodes the fault—the key invariant the cascade relies on. Magnitude-based features are additionally divided by the within-window $|I_1|$ (the load-invariant n_* variants).

Sample counts and class balance. Table V summarises the labelled window counts per dataset and per class. The simulation corpus is balanced across severity (1,680 windows per level, exact to within rounding) and across faulted phase (3,920 per phase); the healthy class is well-represented (6,273 windows, of which 6,138 come from matched-load pre-fault regions of fault runs and 135 from the dedicated healthy runs).

Leakage-safe grouping. All inception-time variants of an operating point and all repetitions of a real-motor condition are bound into a single cross-validation group so that no window from a particular recording can appear in both training and test. The two evaluation protocols then operate on these groups: Stratified Group k -Fold ($k = 5$) for in-distribution estimation, and *Leave-One-Load-Out* (LOLO), where all windows at one load level form the test set in turn—measuring generalization to an operating point never seen in training, the primary metric of this paper (§IV-A).

IV. PERFORMANCE EVALUATION

This section evaluates the proposed framework end-to-end. §IV-A specifies the implementation, splits, and metrics; §IV-B reports the main cascade performance in-distribution and under the primary cross-load (LOLO) protocol; §IV-C analyses how the three tasks differ under LOLO and introduces the load-aware severity fix; §IV-D compares eight model classes on identical splits and tests cross-model significance; §IV-E isolates the contribution of every novel component (load-robustness mechanisms, biomedical features, load-invariant feature selection, and physics-consistent augmentation/tuning) through controlled ablations; §IV-F quantifies probabilistic calibration, noise robustness, and online fault-onset detection; §IV-G validates the cascade within-hardware on a publicly released real-motor dataset; §IV-H reports the measured end-to-end real-time validation on a Jetson Orin Nano edge module; and §IV-I positions the framework relative to the literature and lists the residual limitations.

A. Implementation and Evaluation Protocol

Features are computed in Python (NumPy/Pandas) and models trained with scikit-learn and XGBoost (400–500 trees, depth 5, learning rate 0.05, subsample/colsample 0.8); deep models use PyTorch and TabPFN-2.5 runs on CPU. Two evaluation protocols are used: Stratified Group k -Fold ($k=5$, grouped by operating point) for the in-distribution estimate, and *Leave-One-Load-Out* (LOLO)—all windows at one load form the test set—as the primary metric, measuring generalization to an operating point never seen in training. Reported metrics are accuracy, macro-F1, severity within-one-level accuracy, mean absolute error (MAE) in levels, quadratic-weighted kappa (QWK), and per-phase F1, computed on out-of-fold predictions and reported as a mean over folds. Stages 2–3 are trained on faulty windows but evaluated on the windows the detector predicts faulty, so cascade numbers reflect realistic error propagation; predictions are also aggregated per recording (median severity, majority-vote phase). Scalers are fit on the training partition only, every benchmark model sees identical folds, no per-model tuning is done on test folds, and all randomized components use fixed seeds.

B. Main Cascade Performance: In-Distribution and Cross-Load

Table VI summarizes the deployed cascade. *The central finding is structural rather than numerical:* in-distribution, every stage exceeds 0.97, but only LOLO separates the three tasks into those that genuinely generalize across operating points (detection, macro-F1 0.948; faulted phase, macro-F1 0.997) and the one that does not without intervention (absolute severity). The detection ROC has area 0.994 (Fig. 4) and the faulted-phase confusion matrix (Fig. 5) is near-diagonal with phase C perfect. End-to-end, scoring a window correct when healthy windows are detected healthy and faulty windows have the correct phase and a severity within one level, the gated cascade reaches within-one accuracy 0.949 in-distribution and 0.916 across unseen loads (per-recording 0.927 and 0.918);

TABLE VI
DEPLOYED CASCADE: PER-STAGE AND END-TO-END PERFORMANCE (MEAN OVER FOLDS).

Stage	Metric	Group k -Fold	LOLO (primary)
1 Detection	macro-F1 / ROC-AUC	0.977 / 0.994	0.948
2 Severity (load-aware)	within-1 / QWK	0.997 / 0.987	0.948
3 Phase ID	macro-F1	0.972	0.997
End-to-end	within-1 / per-record	0.949 / 0.927	0.916 / 0.918

TABLE VII
SEVERITY WITHIN-ONE-LEVEL ACCURACY UNDER FEATURE CHOICE. THE LOAD-AWARE MODEL IS THE DEPLOYED CONFIGURATION.

Severity feature set	Group k -Fold	LOLO	No-load
All features (magnitudes incl.)	0.993	0.887	0.39
Load-invariant + $ I_1 $ -norm	0.990	0.729	0.80
+ measured load (proposed)	0.997	0.946	0.95

TABLE VIII
PER-HELD-OUT-LOAD DETECTION (LOLO, MACRO-F1). DIFFICULTY CONCENTRATES AT LIGHT LOADS.

Hold-out load	NL	20%	40%	60%	80%	100%
Detection F1	0.885	0.749	1.000	1.000	1.000	1.000

Fig. 6 contrasts the two protocols. Notably, faulted-phase identification is the only metric for which cross-load *exceeds* in-distribution, a direct consequence of its load-invariance.

C. Cross-Load Analysis and the Load-Aware Severity Fix

SHAP attribution shows detection and phase ride on the *negative-sequence* ratio and angle and on inter-phase RMS differences—all load-invariant ratios—whereas an initial severity model built on absolute magnitudes *collapsed* at the no-load extreme under LOLO (within-one 0.39), because the $|I_2|$ magnitude that encodes severity scales with the operating point. Two complementary remedies recover it (Table VII). First, restricting severity to dimensionless, load-invariant features and $|I_1|$ -normalized magnitudes lifts the no-load case from 0.39 to 0.80. Second—and decisively—supplying the *measured* load (a sensor input, not a label) raises pooled LOLO within-one from 0.73 to 0.95, since most held-out loads are interpolated between neighbours. Severity errors are then almost entirely between adjacent levels (Fig. 7, QWK 0.987), which validates the ordinal treatment. The per-held-out-load breakdown (Table VIII) localizes the residual difficulty to the light-load extremes (no-load and 20%), where a low-severity fault current most overlaps the healthy distribution—a difficulty visible *only* under LOLO.

D. Comparative Model Evaluation

Table IX and Fig. 8 benchmark eight models on *identical* leakage-safe splits, spanning classical feature classifiers, a tabular foundation model, and a sequence deep network. The first observation is that the *in-distribution* task is *saturated*: the six feature classifiers fall within ± 0.01 macro-F1 on detection, so in-distribution accuracy is not a meaningful discriminator

and any ranking on it would be within fold variation. The honest comparison is therefore cross-load. There, three results stand out. (1) The tabular foundation model *TabPFN-2.5* is the strongest detector and uniquely load-robust, scoring macro-F1 0.979 in-distribution *and* 0.979 cross-load, whereas gradient boosting drops from 0.976 to 0.948; it also gives the sharpest in-distribution severity (0.997 within-one, MAE 0.033 level, the lowest of all models) and degrades gracefully across loads (0.969). (2) Severity is the task that separates models: tree ensembles extrapolate poorly to unseen loads (random forest LOLO within-one 0.688), while the MLP and SVM remain robust (0.982, 0.975). (3) Faulted phase is load-robust for nearly every model (LOLO macro-F1 ≥ 0.96), confirming that the negative-sequence-angle signature is model-agnostic. The raw-signal *xLSTM* trails the feature learners on every task (LOLO macro-F1 0.890 detection, within-one 0.908 severity, 0.961 phase) yet improves on the earlier raw-signal deep baselines (1-D ResNet and PCM-Net, 0.898 detection; Table XI), reaffirming the data-efficiency of physically engineered features over raw-waveform deep nets in this regime. The feature pipeline is also lightweight—feature extraction plus XGBoost inference costs ≈ 0.34 ms per window on CPU (Fig. 9), and TabPFN-2.5 needs no task-specific training—so the practitioner can choose foundation-model robustness or gradient-boosted speed without sacrificing accuracy.

Best model per task. Reading the benchmark task-by-task and protocol-by-protocol (Table X), *TabPFN-2.5* is the *only* model that is *simultaneously* top-tier on all three tasks: best (tied) cross-load detector, lowest in-distribution severity error by a wide margin (MAE 0.033 vs. next best 0.095), and within 0.006 macro-F1 of the best phase score. No other model is in the top group on more than two tasks. The complementary observations are that (a) under LOLO severity—the cascade’s hardest task—the MLP edges TabPFN-2.5 (0.982 vs. 0.969 within-one), so for cross-load severity alone the MLP or the FiLM-conditioned PCM-Net (Table XIII) is preferable; (b) for faulted phase under LOLO, the tree ensembles (XGBoost and Random Forest) tie TabPFN-2.5 within 0.006 macro-F1, so any of them suffices; and (c) the raw-signal *xLSTM* is the weakest learner on every cell, confirming the data-efficiency advantage of physical features in this regime. The practical takeaway is that the deployed cascade, currently XGBoost (with the load-aware severity stage), can be upgraded in place by swapping the Stage 1 detector and the in-distribution Stage 2 severity head for TabPFN-2.5, while retaining XGBoost or the MLP for the cross-load Stage 2 operating point—all on the same 72 features and identical splits.

Statistical significance. To test whether these differences are real rather than fold noise we follow the Demšar protocol [22]: a Friedman omnibus test across the seven feature/foundation models on the per-fold (group k -fold, five blocks) and per-load (LOLO, six blocks) primary scores, with Holm-corrected pairwise Wilcoxon signed-rank tests (Table XII). Two conclusions follow. First, *in-distribution detection and phase identification show no significant model differences* ($p = 0.26$ and $p = 0.76$), which quantitatively confirms that the in-distribution task is saturated—the choice among feature models cannot be justified on in-distribution accuracy. Second, the differences that




Fig. 4. Detection. Left: receiver operating characteristic, overall and per held-out load (overall AUC = 0.994). Right: confusion matrix (group k -fold, out-of-fold).

TABLE IX

FULL MODEL BENCHMARK ON IDENTICAL LEAKAGE-SAFE SPLITS: STRATIFIEDGROUPKFOLD (GK, IN-DISTRIBUTION) AND LEAVE-ONE-LOAD-OUT (LOLO, CROSS-LOAD, PRIMARY). DETECTION AND PHASE REPORT MACRO-F1; SEVERITY REPORTS WITHIN-ONE-LEVEL ACCURACY. BEST PER COLUMN IN BOLD. TABPFN-2.5 AND xLSTM ARE ADDED IN THIS WORK.

Model	Detection (macro-F1)		Severity (within-1)		Phase (macro-F1)	
	GK	LOLO	GK	LOLO	GK	LOLO
Logistic Regression	0.978	0.978	0.997	0.971	0.970	0.920
SVM-RBF	0.978	0.977	0.991	0.975	0.984	0.967
k -NN	0.977	0.977	0.937	0.942	0.982	0.981
MLP	0.978	0.978	0.997	0.982	0.956	0.981
Random Forest	0.978	0.979	0.985	0.688	0.965	0.997
XGBoost	0.976	0.948	0.990	0.943	0.976	0.997
TabPFN-2.5	0.979	0.979	0.997	0.969	0.968	0.991
xLSTM	0.957	0.890	0.982	0.908	0.921	0.961




Fig. 5. Faulted-phase confusion (group k -fold, out-of-fold). Phase C is perfect; only minor A $\leftrightarrow$ B confusion remains. Cross-load (LOLO) phase accuracy is 0.997.

matter are significant: absolute severity differs across models under both protocols ($p = 0.0013$ in-distribution, $p = 0.0008$ cross-load) and detection differs across loads ($p = 0.050$), driven by the gradient-boosted model's cross-load drop and




Fig. 6. Per-stage and end-to-end performance, in-distribution (group k -fold) vs. cross-load (LOLO). Phase identification is the only metric for which cross-load exceeds in-distribution, reflecting its load-invariance.

the tree ensembles' poor severity extrapolation (Random Forest LOLO within-one 0.70). By average Friedman rank, TabPFN-2.5 and the MLP are strongest on cross-load severity and the tree ensembles weakest. With only five-to-six blocks, however, no *pairwise* comparison survives Holm correction (smallest adjusted $p = 0.19$); the omnibus is therefore the reliable signal, and localizing pairwise differences would require more folds or additional operating points—a limitation we flag rather than




Fig. 7. Severity confusion (group k -fold). Errors are almost entirely between adjacent levels, justifying the ordinal treatment and the within-one-level metric (QWK = 0.987).

TABLE X

BEST MODEL PER DIAGNOSTIC TASK AND PROTOCOL, DISTILLED FROM TABLE IX. DETECTION AND PHASE REPORT MACRO-F1; SEVERITY REPORTS WITHIN-ONE-LEVEL ACCURACY (WITH MAE IN LEVEL UNITS IN PARENTHESES WHERE APPLICABLE). “TIED WITHIN 0.001” IS ROUNDED.

Task	Protocol	Best model	Score
Detection	GroupKFold	TabPFN-2.5	0.979
Detection	LOLO	TabPFN-2.5 / RF (tied)	0.979
Severity	GroupKFold	TabPFN-2.5	0.997 (MAE 0.033)
Severity	LOLO	MLP	0.982 (MAE 0.112)
Phase	GroupKFold	SVM-RBF	0.984
Phase	LOLO	XGBoost / RF (tied)	0.997

Overall (top-tier on all three tasks under both protocols)

TabPFN-2.5—best or tied-best on 4/6 cells, within 0.013 of best on the other 2.

TABLE XI

FEATURE-BASED CASCADE VS. A RAW-SIGNAL DEEP BASELINE.

Task / metric	XGBoost (features)	ResNet-1D (raw)
Detection F1 (group k -fold)	0.977	0.898
Detection F1 (LOLO, load 20%)	0.749	0.401
Phase macro-F1 (group k -fold)	0.972	0.848
Severity within-1 (group k -fold)	0.990	1.000

TABLE XII

SIGNIFICANCE ACROSS THE SEVEN FEATURE/FOUNDATION MODELS (FRIEDMAN OMNIBUS ON PER-FOLD / PER-LOAD PRIMARY SCORES; DETECTION AND PHASE USE MACRO-F1, SEVERITY WITHIN-ONE-LEVEL ACCURACY).

Task	Protocol	χ^2	p	Outcome
Detection	group k -fold	7.66	0.264	n.s.
Detection	LOLO	12.61	0.050	sig.
Severity	group k -fold	21.84	0.0013	sig.
Severity	LOLO	22.85	0.0008	sig.
Phase	group k -fold	3.35	0.764	n.s.
Phase	LOLO	6.00	0.423	n.s.

over-claim from.

TABLE XIII
LOAD-ROBUSTNESS MECHANISMS (LOLO).

Mechanism	Cross-load result
SCST representation	severity within-1 0.75 $\rightarrow$ 1.0; phase needs angle
LIR-mRMR selection	$k=10$ severity within-1 0.740 $\rightarrow$ 0.848; drift halved
PCM-Net + FiLM	severity within-1 0.932 $\rightarrow$ 0.982 (best; > load-aware GBM 0.946)
PADA (sim $\rightarrow$ real)	detection F1 0.48 $\rightarrow$ 0.98 (few-shot); 0.50 phase (unsup.)
PC-Diff	current-balance exact; generation not controllable (negative)

E. Mechanism Ablations: Load-Robustness Components

Table XIII ablates the four proposed mechanisms under LOLO, the protocol on which each is designed to help. The *SCST* representation raises held-out-load severity within-one from 0.75 (raw-phase Stockwell) to 1.0; consistent with the physics, its magnitude tensor cannot localize phase, which is why phase uses the negative-sequence angle instead. *LIR-mRMR* beats plain mRMR at compact size ($k=10$ within-one 0.848 vs. 0.740) with roughly half the cross-load feature drift, yielding a sparse, interpretable subset. *PCM-Net*’s FiLM load-conditioning lifts LOLO severity within-one from 0.932 to **0.982**—the best in the study, exceeding even the load-aware gradient-boosted model (0.946); FiLM does not help the load-invariant phase task, exactly as the physics predicts, which is itself evidence the mechanism does what it claims. *PADA* closes the simulation-to-experiment gap on the public real-motor dataset: a naively transferred detector scores macro-F1 0.48 (it collapses to predicting all-faulty under covariate shift), physics-anchored unsupervised alignment reaches 0.50 on the hard phase task, and few-shot calibration with two real repetitions recovers detection to **0.98**. Finally, the physics-constrained generator enforces current balance exactly (residual 0.000 vs. 0.991 for an unconstrained generator) but did not yield severity-controllable samples under our compute budget—an honest negative result that we report to motivate a stronger conditional generator rather than omit.

F. Biomedical Anomaly Features

Table XIV isolates the contribution of the biomedical anomaly block (Section II-F). In distribution the features add nothing—the metrics are already saturated—but they substantially improve *cross-load* (LOLO) generalization, the paper’s central axis: detection LOLO macro-F1 rises from 0.948 to 0.978 (+0.030) and severity within-one-level accuracy from 0.942 to 0.979 (+0.037), while the already load-robust phase task is unchanged. By gradient-boosting gain the most useful biomedical features are the Hjorth mobility/complexity of the Park modulus, the Higuchi/Petrosian fractal dimension and permutation entropy of the residual, and dispersion entropy—all amplitude-scale-free, which is why they transfer across loads better than absolute-magnitude physics features and directly attack the cross-load severity bottleneck. This cross-domain import of biomedical complexity/irregularity features into stator-current diagnosis is, to our knowledge, new.




Fig. 8. Model benchmark across the three diagnostic tasks under both protocols (GroupKFold vs. LOLO). Hatched bars denote the two models added in this work (TabPFN-2.5, xLSTM). Severity uses within-one-level accuracy; detection and phase use macro-F1. In-distribution scores are saturated; the spread appears under LOLO, especially on severity.




Fig. 9. Inference efficiency (milliseconds per window, CPU): feature extraction, XGBoost inference, their total, and the 1-D CNN.

TABLE XIV

BIOMEDICAL-FEATURE ABLATION (XGBOOST, IDENTICAL LEAKAGE-SAFE SPLITS):
PHYSICS FEATURES VS. PHYSICS + BIOMEDICAL ANOMALY FEATURES.
DETECTION/PHASE: MACRO-F1; SEVERITY: WITHIN-ONE-LEVEL ACCURACY.

Task (metric)	Protocol	Physics	Physics+Bio	Δ
Detection (macro-F1)	10-fold GK	0.977	0.977	+0.000
Detection (macro-F1)	LOLO	0.948	0.978	+0.030
Severity (within-1)	10-fold GK	0.997	0.997	+0.000
Severity (within-1)	LOLO	0.942	0.979	+0.037
Phase (macro-F1)	10-fold GK	0.964	0.964	+0.000
Phase (macro-F1)	LOLO	0.997	0.997	+0.000

G. Feature Importance and Load-Invariant Selection

Importance. Fig. 10 ranks the features that drive the models. For detection, SHAP attribution (a) is led by the Park-vector major axis and eccentricity, the zero/negative-sequence ratios, inter-phase RMS differences, and the EPVA severity factor—all physical inter-turn signatures, confirming the models exploit fault physics rather than dataset shortcuts. Among the novel biomedical features (b), the Park-modulus Hjorth parameters, the residual Higuchi/Petrosian fractal dimension

and permutation entropy, and dispersion entropy carry the most weight on the (hardest) severity task, explaining their cross-load gains (Table XIV). *Reduction.* Because many features are redundant or load-dependent, we reduce the bank with the load-invariance-regularized criterion of Section II-F (LIR-mRMR). Fig. 11 compares it with plain mRMR: at a compact $k=10$ features LIR-mRMR reaches cross-load (LOLO) severity within-one of 0.848 versus mRMR’s 0.740 (a) while roughly halving the selected features’ across-load drift (b, 0.18 vs 0.38), because it explicitly penalizes features whose class-conditional mean shifts with load. The selected subset is dimensionless and physics-traceable (negative-sequence angle $\sin/\cos$, Park eccentricity, $|I_1|$ -normalized imbalance, per-phase skew)—a desirable property for safety-critical monitoring—and approaches the full 72-feature accuracy with an order of magnitude fewer inputs. The $k=15$ dip reflects greedy-selection noise, not a benefit of additional features.

H. Hyperparameter Tuning and Physics-Consistent Augmentation

We finally isolate the effect of leakage-safe hyperparameter tuning and physics-consistent training-fold augmentation on the deployed XGBoost pipeline, evaluated under *10-fold* StratifiedGroupKFold (the 147 operating-point groups satisfy $n_{\text{folds}} \leq n_{\text{groups}}$) and LOLO. Hyperparameters are tuned for cross-load generalization (Section II-F) and frozen for both protocols; augmentation is applied to the training fold only and re-featurised. Table XV reports the three conditions. Tuning alone is marginal on the already-saturated metrics (consistent with the significance analysis of Section IV-D), but *augmentation is the decisive lever for cross-load generalization*: it raises severity within-one-level accuracy from 0.911 to 0.935 under LOLO and from 0.934 to 0.976 in-distribution, cutting the mean absolute error by 20% ($0.917 \rightarrow 0.731$ level), and lifts detection LOLO macro-F1 from 0.964 to 0.977, while faulted-phase identification remains at 0.997. This confirms the augmentation principle: synthesising training windows




Fig. 10. Feature importance. (a) Stage-1 detection by mean |SHAP| (top-12): Park-vector geometry, sequence-component ratios, inter-phase RMS differences, and the EPVA factor dominate—all physical inter-turn signatures. (b) Most influential *biomedical* features (gradient-boosting gain) on the severity task: Park-modulus Hjorth parameters, residual fractal dimension and permutation entropy, dispersion entropy.




Fig. 11. Load-invariance-regularized feature reduction (LIR-mRMR) vs. plain mRMR, vs. feature-set size k . (a) Cross-load (LOLO) severity within-one accuracy—LIR-mRMR is higher at compact sizes. (b) Mean across-load drift of the selected features (lower is better)—LIR-mRMR roughly halves it by penalising load-unstable features.

TABLE XV

EFFECT OF LEAKAGE-SAFE TUNING AND PHYSICS-CONSISTENT TRAINING-FOLD AUGMENTATION ON THE DEPLOYED XGBOOST PIPELINE (10-FOLD STRATIFIEDGROUPKFOLD AND LOLO). DETECTION/PHASE: MACRO-F1; SEVERITY: WITHIN-ONE-LEVEL ACCURACY. BEST PER ROW IN BOLD.

Task (metric)	Protocol	Default	Tuned	Tuned+Aug
Detection (macro-F1)	10-fold GK	0.976	0.976	0.979
Detection (macro-F1)	LOLO	0.964	0.964	0.977
Severity (within-1)	10-fold GK	0.934	0.937	0.976
Severity (within-1)	LOLO	0.911	0.908	0.935
Phase (macro-F1)	10-fold GK	0.964	0.936	0.964
Phase (macro-F1)	LOLO	0.994	0.997	0.997

toward unseen operating points—by jointly scaling the three phases, so the inter-phase imbalance and current balance are preserved—directly improves cross-load diagnosis without fabricating the fault signature.

I. Robustness, Calibration, and Onset Detection

Stage-1 probabilities are well calibrated ($ECE = 0.021$, $Brier = 0.021$), so the detector’s confidence can gate maintenance actions directly. Conformalized quantile regression improves severity-interval coverage over naive quantiles, with group-conditional calibration noted as the remedy for residual under-coverage. The onset detector (Table XVII) flags 100% of faults with no pre-fault false alarms; for severity $\geq 0.5\%$ the fault is caught within the first faulted cycle, and only the smallest 0.3% faults incur a lag. Under additive noise (Table XVI), severity and phase are essentially flat to 20 dB SNR while detection is the noise-sensitive stage—expected, since separating a small fault from healthy operation is hardest in noise. Training is stable for both the gradient-boosted and deep models (Fig. 12), and the per-stage precision/recall/F1 breakdown is given in Fig. 13.




Fig. 12. Training behaviour. Left: XGBoost detection convergence (train/validation log-loss vs. boosting round). Right: 1-D CNN learning curves (training loss and validation macro-F1 vs. epoch), showing stable convergence without over-fitting.

TABLE XVI
NOISE ROBUSTNESS (TRAIN CLEAN, TEST NOISY; GROUP k -FOLD).

SNR	Detect macro-F1	Severity within-1	Phase acc.
clean	0.977	0.986	0.958
40 dB	0.927	0.980	0.954
30 dB	0.854	0.975	0.945
20 dB	0.792	0.971	0.939

TABLE XVII
FAULT-ONSET DETECTION (SAMPLED CASES).

Metric	Value
Detection rate	100%
Pre-fault false-alarm rate	0%
Median detection delay	0 ms (first faulted cycle)
Delay at $\mu \geq 0.5\%$ / $\mu = 0.3\%$	0 ms / ≈ 1.4 s




Fig. 13. Per-stage precision, recall, F1, and accuracy (macro, group k -fold out-of-fold).

J. Within-Hardware Validation on a Real Motor

All results so far are from simulation; we now train and test the same models *on a real motor*. Using the public ibarram dataset [35] (a 0.75 hp IM, three-phase current at 1 kHz, no load; 65 recordings = 13 classes $\times$ 5 repetitions), we evaluate every model within-hardware under two leakage-safe protocols: StratifiedGroupKFold grouped by recording, and Leave-One-Repetition-Out (LORO), which holds out an entire unseen recording per class. Severity uses the dataset's

own four-level scale (10/20/30/40%). Table XVIII reports the outcome and Fig. 14 the XGBoost confusion matrices, and we deliberately report the weaknesses as fully as the strengths.

Three findings stand out. First, *detection transfers to hardware*: TabPFN-2.5 and the gradient-boosted/forest models reach macro-F1 0.95–0.99 on the real motor, and the LORO score (an entirely unseen recording) is as strong as the in-domain one, so fault *presence* is decided reliably from real current. Second, *severity is moderate*: within-one-level accuracy is 0.83–0.96 but exact accuracy falls to 0.6–0.75 (TabPFN-2.5 best, with the lowest mean absolute error of 0.34 level), reflecting genuinely harder discrimination of adjacent shorted-turn fractions on real signals. Third, *per-window* faulted-phase localization is weak on hardware (macro-F1 0.64–0.72, versus 0.97 in simulation), because the per-window negative-sequence *angle* estimate is noisy at the low 1 kHz rate. However, the realistic unit of diagnosis is a recording, not a single window: *aggregating a recording's windows by majority vote recovers phase macro-F1 to 0.85 (group k -fold) and 0.93–0.95 (LORO)* for the gradient-boosted and forest models—a +0.22 gain under LORO that requires no new features. A controlled ablation found that reference-invariant feature engineering (relative angle, per-phase asymmetry) did *not* improve beyond this, which localizes the cause to per-window estimator variance rather than feature design; closing the residual per-window gap (higher sampling, supervised phase calibration on the target machine) is left to future work. Consistent with the simulation benchmark, TabPFN-2.5 is again the strongest tabular model, and the raw-signal xLSTM trails on detection and phase (its high severity within-one reflects adjacent-level predictions rather than exact accuracy).

K. Real-Time Edge Deployment on a Jetson Orin Nano

Why XGBoost for deployment. The eight-model benchmark (§IV-D) showed several classifiers within fold noise of the best on the primary cross-load metric—TabPFN-2.5 holds a 0.002 edge on detection LORO macro-F1 and the MLP a 0.04 edge on cross-load severity within-one accuracy, with all top models clustered above 0.97 in-distribution. Selecting




Fig. 14. Within-hardware confusion matrices on the real motor (XGBoost, GroupKFold out-of-fold): detection (left), four-level severity (middle), faulted phase (right). Detection is clean and severity errors are near-diagonal, but the phase matrix shows substantial off-diagonal confusion—the main simulation-to-reality gap.

TABLE XVIII

WITHIN-HARDWARE RESULTS ON THE REAL IBARRAM MOTOR (STRATIFIEDGROUPKFOLD BY RECORDING, GK; AND LEAVE-ONE-REPETITION-OUT, LORO). DETECTION AND PHASE REPORT MACRO-F1; SEVERITY REPORTS WITHIN-ONE-LEVEL ACCURACY ON ITS FOUR-LEVEL SCALE. BEST PER COLUMN IN BOLD. PHASE VALUES ARE PER-WINDOW; PER-RECORDING MAJORITY VOTE RAISES XGBOOST PHASE MACRO-F1 TO 0.85 (GK) AND 0.95 (LORO).

Model	Detection		Severity (w-1)		Phase	
	GK	LORO	GK	LORO	GK	LORO
Logistic Regression	0.817	0.864	0.830	0.818	0.664	0.678
SVM-RBF	0.877	0.923	0.866	0.818	0.661	0.689
k-NN	0.829	0.884	0.827	0.804	0.642	0.648
MLP	0.846	0.880	0.866	0.811	0.664	0.690
Random Forest	0.923	0.947	0.900	0.829	0.650	0.718
XGBoost	0.966	0.955	0.900	0.836	0.636	0.707
TabPFN-2.5	0.994	0.954	0.929	0.864	0.643	0.701
xLSTM	0.708	0.699	0.963	0.957	0.651	0.648

the *deployed* model is therefore not an accuracy decision but a deployment-trade-off decision. Table XIX summarises the Pareto comparison across the axes that actually matter at the motor edge. XGBoost is uniquely top-tier on all three diagnostic tasks while also being the smallest bundle (5 MB of JSON), the fastest to load and predict (sub-millisecond per booster after the Booster+DMatrix switch of §IV-A), the lowest in runtime dependencies (a single pip wheel), and the most directly auditable (SHAP attributions on engineered physical features, §IV-G). The two competing models lose the trade-off on different axes: TabPFN-2.5 [12] is a ≈ 500 MB gated transformer whose in-context inference fits the entire training set into the prompt at each restart (minutes-scale warm-up), and the MLP requires a deep-learning runtime for a severity gain that the load-aware XGBoost cascade (§IV-C) already closes to within 0.005. The deployment claim of this paper is therefore not “XGBoost is the best classifier” but “XGBoost is the Pareto-optimal classifier for the operational constraints of an on-motor edge node”—an explicit, evidence-backed deployment rationale.

To verify that the deployed cascade is actually deployable, we built an end-to-end real-time streaming pipeline and ran it on physical NVIDIA Jetson hardware: a host laptop transmits

TABLE XIX

PARETO COMPARISON OF THE THREE STRONGEST CANDIDATE MODELS FOR THE EDGE DEPLOYMENT. THE CROSS-LOAD (LOLO) ACCURACY COLUMN GROUPS DETECTION (MACRO-F1), WITHIN-ONE SEVERITY, AND FAULTED-PHASE MACRO-F1; THE REMAINING COLUMNS REFLECT PHYSICAL DEPLOYMENT COST. BEST PER COLUMN IN BOLD; “ $\approx$ ” MARKS WITHIN-FOLD-NOISE DIFFERENCES.

Model	LOLO acc. (D/S/P)	Bundle	Predict	Cold-start
XGBoost (deployed)	0.98 / 0.95 / 1.00[†]	5 MB	1 ms[‡]	<0.1 s
TabPFN-2.5	$\approx 0.98 / 0.97 / 0.99$	~ 500 MB	~ 50 – 500 ms	minutes
MLP	$\approx 0.98 / 0.98 / 0.98$	~ 5 MB	few ms	~ 1 s

[†] XGBoost detection is reported with the tuning + augmentation pipeline of §IV-H; severity

[‡] Measured on the Jetson Orin Nano with the native Booster+DMatrix API (§IV-A); thro

the three-phase stator current of a real motor to an NVIDIA Jetson Orin Nano (Super, JetPack 6, R36.4.7) over a USB type-C / CDC-ACM serial channel, and the Jetson buffers, windows, extracts the physics and biomedical features, and runs the gated cascade (onset $\rightarrow$ detection $\rightarrow$ {severity, phase}) at the acquisition rate.

System architecture. The host (laptop) acts as a programmable current source: it streams the three-phase samples of a chosen recording over a binary UART/CDC-ACM channel at 921,600 baud, packed as 14-byte frames (0xAA sync + six 16-bit Q15 big-endian channels + 0x55 checksum) and paced at a configurable wire rate (5 kHz here, an embedded-ADC-realistic rate, upsampled at the sender from the recording’s native 1 kHz). The Jetson runs a polling reader on /dev/ttyGS0 that (i) resyncs on the framing bytes and decodes the six channels (the three voltage channels are discarded; the three current channels are kept), (ii) resamples the wire stream from 5 kHz back to the model’s training rate of 1 kHz with a polyphase FIR (scipy.signal.resample_poly), (iii) accumulates samples into a circular buffer and emits a $W=500$ -sample window every stride=250 samples, (iv) computes the 72 physics features and 24 biomedical features per window, and (v) invokes the deployed XGBoost cascade, printing per-window predictions and a rolling per-recording verdict (majority-vote phase, median severity, healthy/faulty by window majority). The deployed model bundle (detect.json, severity.json, phase.json, meta.json) is trained off-

TABLE XX

PER-WINDOW LATENCY OF THE DEPLOYED CASCADE ON THE JETSON ORIN NANO (SUPER) UNDER THREE EXECUTION MODES, AVERAGED OVER 200 INFERENCES AFTER 5 WARM-UP CALLS. FEATURE EXTRACTION IS IDENTICAL ACROSS MODES (NumPy/SciPy on CPU).

Mode	Feature (ms)	Predict (ms)	Total (ms)
CPU (sklearn wrapper)	22.7	3.5	26.2
GPU (XGBoost device='cuda')	—	—	requires source build
Accelerated CPU (Booster + DMatrix)	22.7	1.0	23.7

Window-fill at $W=500$, $f_s=1$ kHz: 500 ms (12–20× inference headroom).
Decision cadence: 4 Hz (every stride / $f_s = 250/1000$ s).

line on the host with leakage-safe splits, exported, and copied once to the Jetson; no retraining occurs at the edge. The same protocol works with a live ADC—the sender is simply replaced by an acquisition daemon—so the wire-protocol loopback validates the entire software path that would carry a real motor’s current to the diagnosis.

Measured per-window latency (CPU vs. GPU vs. accelerated CPU). We measured the end-to-end per-window cost of the deployed cascade on the Jetson Orin Nano under three execution modes (Table XX). *CPU*: the production path—NumPy/SciPy feature extraction on the ARM Cortex-A78 cluster and three XGBoost predictors via the scikit-learn wrapper—averages 22.7 ms for the 96-feature extraction and 3.5 ms for the three booster calls, totalling 26.2 ms per window over 200 repetitions. *GPU*: the on-board Orin GPU (Ampere, CUDA 12.6 reported by `nvidia-smi`) accepts `device='cuda'` in XGBoost 3.2 but the pre-built aarch64 pip wheel does not ship CUDA kernels compiled for the Jetson Orin compute capability (`sm_87`)—first prediction returns `cudaErrorNoKernelImageForDevice`. Operationally this is recoverable by a source build of XGBoost with `-DGPU_COMPUTE_CAPABILITIES=87`, but for the deployed tree cascade the GPU is unlikely to win: at 1 row × 96 features per prediction the host-to-device transfer dominates the actual tree traversal, and we therefore report it as a present-but-unused option whose path would only pay off for the deep-model alternative (Section II-G), where the GPU+TensorRT pipeline is intended to handle PCM-Net / xLSTM streaming inference. *Accelerated CPU*: switching from the scikit-learn wrapper to XGBoost’s native `Booster + DMatrix` C-level API drops the three predictions from 3.5 ms to 1.0 ms (3.5× over the production path) and brings the end-to-end total to 23.7 ms; we therefore recommend the native Booster API for shipped binaries. In all cases the inference cost is dwarfed by the 500 ms window-fill latency at $W=500$, $f_s=1$ kHz: the pipeline issues a decision every 250 ms (stride / f_s) and the inference budget has roughly 12× headroom over the worst (CPU) mode. The onset trigger (Stage 0) runs at one decision per fundamental cycle and is essentially free.

Measured validation on physical Jetson hardware. Six held-out repetition-05 recordings of the public real-motor dataset—one healthy (HLT) and five faulted recordings spanning all three phases at severities {10, 30, 40}%—were streamed end-to-end from the host laptop to the Jetson Orin Nano. Each 5 s recording was upsampled 1 kHz → 5 kHz at the sender, Q15-

TABLE XXI

REAL-TIME EDGE CASCADE ON A JETSON ORIN NANO (SUPER): PER-RECORDING VERDICTS ON SIX HELD-OUT REPETITION-05 RECORDINGS STREAMED OVER A USB TYPE-C / CDC-ACM CHANNEL AT 5 kHz WIRE RATE (921,600 BAUD, 0 BYTES DROPPED OVER 30 s). VERDICTS ARE THE ROLLING 10-WINDOW AGGREGATE (MAJORITY-VOTE PHASE, MEDIAN SEVERITY, HEALTHY/FAULTY BY WINDOW MAJORITY). THE HEALTHY HLT CASE IS THE FALSE POSITIVE; PHASE ID IS EXACT ON EVERY FAULTED RECORDING.

Case (rep 05)	Ground truth	Jetson verdict	Outcome
HLT	Healthy	FAULTY / 40% / A	FP (detect)
B30	Faulty / 30% / B	FAULTY / 30% / B	all correct
A30	Faulty / 30% / A	FAULTY / 40% / A	phase / within-1
C30	Faulty / 30% / C	FAULTY / 40% / C	phase / within-1
A40	Faulty / 40% / A	FAULTY / 40% / A	all correct
B10	Faulty / 10% / B	FAULTY / 40% / B	phase only

Aggregate per-recording: detect 5/6, phase 5/5 faulted, severity within-1 4/5.
Wire: ≈ 4880 pkt/s, 0.55 Mbps, 0 bytes dropped over 30 s.

Per-window latency: see Table XX.

packed into the 14-byte frames, and transmitted over the CDC-ACM channel at ≈ 4880 packets s^{-1} ; the receiver dropped zero bytes across the 30 s sweep. Fig. 15 visualizes the per-recording verdicts side-by-side with ground truth (green = correct, yellow = within one severity level, red = wrong) and the per-window latency under the three execution modes (CPU, GPU, accelerated CPU). Table XXI reports the same verdicts in tabular form. Concretely: (i) detection succeeded on 5/6 recordings, the HLT healthy case being the lone false positive at the incipient-detection boundary already documented in Section IV-J—a known weakness of the deployed bundle on real-motor data, traceable to the absent healthy baseline in a single-window stream; (ii) faulted-phase identification succeeded on 5/5 faulted recordings (100% accuracy, matching the per-recording phase macro-F1 0.93–0.95 measured in Section IV-J); and (iii) severity was exact on 2/5 cases (B30 and A40) and within one severity level on 4/5, with the only large miss being the most incipient B10 (10%) case over-predicted as 40%, mirroring the within-one accuracy 0.83–0.96 reported for this bundle in Section IV-J. *Operational interpretation.* The deployed pipeline therefore correctly raises an alarm for every actual fault, names the right phase every time, and supplies a usable severity reading—the false-positive on healthy and the over-prediction of severity for the 10% case are the two limitations a field operator would experience, both already characterized in Section IV-J.

The deployed bundle deliberately excludes the compensation-baseline-dependent I_2^c features so that the streamed window is fully self-contained; the equivalent uncompensated $|I_2|/|I_1|$ family carries the same signal. For a different motor or sensing chain, only the recording’s (f_s, f_0) pair needs to be re-estimated (a short healthy pre-roll suffices); the feature extractor is already parameterized in (f_s, f_0) and the model can be re-fitted on the new rig with the same off-line pipeline. This closes the loop from the off-line evaluation of Sections IV-A–IV-H to a runnable edge deployment, and supports the deployment claim of the Discussion.

Real-Time Edge Validation on a Jetson Orin Nano (Super) — 5 kHz UART, 0 bytes dropped

Per-recording verdicts from the Jetson Orin Nano
(green = correct, yellow = within 1 severity level, red = wrong)




Fig. 15. End-to-end real-time validation of the deployed cascade on a Jetson Orin Nano (Super, JetPack 6, R36.4.7) over a USB type-C / CDC-ACM serial channel at 5 kHz wire rate. *Top*: per-recording verdicts vs. ground truth for six held-out repetition-05 recordings (1 healthy, 5 faulted spanning every phase and severities {10, 30, 40}%). Green = correct, yellow = within 1 severity level, red = wrong. Detection 5/6 (only HLT false-positive), faulted phase 5/5, severity within-1 4/5. *Bottom-left*: aggregate per-recording accuracy. *Bottom-right*: measured per-window latency on the Jetson Orin Nano under three execution modes (CPU sklearn, GPU XGBoost device='cuda' which requires a source build for sm_87, and accelerated CPU via XGBoost native Booster + DMatrix). Dotted line is the 500 ms window-fill at $f_s=1$ kHz, $W=500$.

L. Discussion: Positioning, Limitations, and Future Work

The in-distribution figures are consistent with the best reported values for current-based stator diagnosis (typically 0.96–1.00 accuracy [3], [8], [9]); the contribution here is not a higher in-distribution number—which the saturated benchmark shows is no longer informative—but a demonstration of *which* of those numbers survive a change of operating point, and a set of mechanisms that make the load-dependent ones survive. The in-distribution-versus-LOLO severity gap (within-one 0.99 vs. 0.39 before the load-aware fix) quantifies how much a conventional evaluation would have over-promised; we therefore advocate cross-condition (LOLO) reporting, alongside interpretability evidence, as standard practice. From a deployment standpoint the pipeline is attractive: it consumes only the already-measured three-phase current, exposes interpretable physical quantities (the negative-sequence ratio and EPVA factor that drive a decision), and provides calibrated probabilities with an online onset trigger suitable for a protection relay. *Model choice for the edge node is an explicit Pareto decision.* Because the eight-model benchmark (§IV-D) shows the top tabular learners clustered within ± 0.005 macro-F1 on cross-load detection and within ± 0.04 within-one accuracy on cross-load severity (with the load-aware XGBoost cascade of §IV-C closing the latter gap to ± 0.005), the choice of *deployed* model is dictated by deployment cost rather than accuracy.

The Pareto comparison of Table XIX and the measured Jetson Orin Nano latency of Table XX together justify the gradient-boosted cascade as the shipped solution: a 5 MB JSON bundle, three boosters totalling 1 ms per prediction via the native API, a single xgboost runtime dependency, sub-second cold-start, and SHAP-auditable physical features—each axis on which the foundation-model option (TabPFN-2.5: ~ 500 MB gated transformer, in-context fit at every restart) or the deep-learning alternatives (PCM-Net / xLSTM: PyTorch runtime, source-built CUDA kernels on the Jetson) lose decisively. TabPFN-2.5 remains the recommended choice when the deployment target is a back-end inference server rather than a motor-mounted edge node. Within-hardware validation on a real motor (Section IV-J) confirms that detection transfers (macro-F1 up to 0.99) and severity is usable (within-one ≥ 0.86); per-window faulted-phase localization is weak (macro-F1 0.72) but per-recording aggregation—the realistic diagnostic unit—recovers it to 0.93–0.95. The remaining limitations are that the load-robustness results are validated only in simulation (the public real dataset has no load variation), and that a full instrumented campaign with supply-voltage unbalance and richer real severities remains future work.

V. CONCLUSION

We presented a physics-informed, leakage-safe machine-learning framework for stator inter-turn fault diagnosis in induction motors that performs onset detection, healthy/faulty detection, ordinal severity estimation, and faulted-phase identification from the three-phase stator current. By adopting Leave-One-Load-Out evaluation we showed that detection and faulted-phase identification are intrinsically load-robust, while absolute severity is load-dependent and is recovered to within-one accuracy of 0.95 across unseen loads once the model is made load-aware. A controlled eight-model benchmark—including the tabular foundation model TabPFN-2.5 and an xLSTM sequence network—demonstrated that the in-distribution task is saturated and that cross-load behaviour is the meaningful axis, on which TabPFN-2.5 is the most load-robust detector. Four mechanisms (a sequence-component Stockwell tensor, a load-invariance-regularized feature selection, a FiLM-conditioned multi-task network, and physics-anchored domain adaptation) each independently improved cross-load performance, and the framework was validated simulation-to-experiment on a public real-motor dataset. The end-to-end cascade attains within-one accuracy of 0.95 in-distribution and 0.92 across unseen loads, with calibrated detection, graceful noise degradation, and reliable onset detection. A within-hardware evaluation of all eight models on the real motor confirmed that detection transfers (macro-F1 up to 0.99) and severity is usable; per-window faulted-phase localization is the main gap (macro-F1 0.72) but per-recording aggregation recovers it to 0.93–0.95. Future work includes closing the residual per-window phase gap (higher sampling, supervised phase calibration), a stronger conditional generator for physics-constrained augmentation (the signature-preserving training-fold augmentation introduced here already cut cross-load severity error by 20%), full selective-state-space and Kolmogorov–Arnold backbones, an experimental campaign with supply voltage-unbalance rejection, and group-conditional conformal calibration of the severity intervals. All code, dataset splits, and result artifacts are released for reproducibility.

REFERENCES

- [1] J. Penman, H. G. Sedding, B. A. Lloyd, and W. T. Fink, "Detection of inter-turn short circuits in the stator windings of operating motors," *IEEE Trans. Energy Convers.*, vol. 9, no. 4, pp. 652–658, 1994.
- [2] A. J. M. Cardoso, S. M. A. Cruz, and D. S. B. Fonseca, "Inter-turn stator winding fault diagnosis in three-phase induction motors, by Park's vector approach," *IEEE Trans. Energy Convers.*, vol. 14, no. 3, pp. 595–598, 1999.
- [3] R. Hussain, M. Alshaikh Saleh, and S. S. Refaat, "Various faults classification of industrial application of induction motors using supervised machine learning: A comprehensive review," *IEEE Access*, vol. 13, pp. 146649–146675, 2025.
- [4] D. García-Pérez *et al.*, "Machine learning for inverter-fed motors monitoring and fault detection: An overview," *IEEE Access*, vol. 12, 2024.
- [5] S. Ruzimov *et al.*, "Detection of inter-turn short-circuit faults for inverter-fed induction motors based on negative-sequence current analysis," *Sensors*, vol. 25, no. 15, art. 4844, 2025.
- [6] M. A. Mazzoletti, G. R. Bossio, P. D. Donolo, and M. Donolo, "Turn-to-turn short-circuit fault and high-resistance connection diagnosis for induction motor during start-up transient," *IEEE Trans. Ind. Appl.*, vol. 61, no. 2, pp. 1927–1935, 2025.
- [7] M. Singh and A. G. Shaik, "Incipient fault detection in stator windings of an induction motor using Stockwell transform and SVM," *IEEE Trans. Instrum. Meas.*, vol. 69, no. 12, pp. 9496–9504, 2020.
- [8] P. Pietrzak and M. Wolkiewicz, "Machine learning-based stator current data-driven PMSM stator winding fault diagnosis," *Sensors*, vol. 22, no. 24, art. 9668, 2022.
- [9] T. Aasar *et al.*, "Predicting stator winding short-circuit faults in induction motors using machine learning: A comparative study," *Discover Computing*, vol. 29, art. 103, 2026.
- [10] U. Ali, U. Ramzan, W. Ali, and K. A. Al-Jaafari, "An improved fault diagnosis strategy for induction motors using weighted probability ensemble deep learning," *IEEE Access*, vol. 13, pp. 106958–106973, 2025.
- [11] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. ACM SIGKDD*, 2016, pp. 785–794.
- [12] N. Hollmann *et al.*, "Accurate predictions on small data with a tabular foundation model," *Nature*, vol. 637, pp. 319–326, 2025.
- [13] M. Beck *et al.*, "xLSTM: Extended long short-term memory," in *Proc. NeurIPS*, 2024.
- [14] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in *Proc. NeurIPS*, 2017, pp. 4765–4774.
- [15] Y. Romano, E. Patterson, and E. Candès, "Conformalized quantile regression," in *Proc. NeurIPS*, 2019.
- [16] A. N. Angelopoulos and S. Bates, "A gentle introduction to conformal prediction and distribution-free uncertainty quantification," *arXiv:2107.07511*, 2021.
- [17] C. Guo, G. Pleiss, Y. Sun, and K. Q. Weinberger, "On calibration of modern neural networks," in *Proc. ICML*, 2017, pp. 1321–1330.
- [18] T. Ince, S. Kiranyaz, L. Eren, M. Askar, and M. Gabbouj, "Real-time motor fault detection by 1-D convolutional neural networks," *IEEE Trans. Ind. Electron.*, vol. 63, no. 11, pp. 7067–7075, 2016.
- [19] K. He, X. Zhang, S. Ren, and J. Sun, "Deep residual learning for image recognition," in *Proc. IEEE CVPR*, 2016, pp. 770–778.
- [20] A. Gu and T. Dao, "Mamba: Linear-time sequence modeling with selective state spaces," *arXiv:2312.00752*, 2023.
- [21] Z. Liu *et al.*, "KAN: Kolmogorov–Arnold networks," in *Proc. ICLR*, 2025.
- [22] J. Demšar, "Statistical comparisons of classifiers over multiple data sets," *J. Mach. Learn. Res.*, vol. 7, pp. 1–30, 2006.
- [23] I. Gulrajani and D. Lopez-Paz, "In search of lost domain generalization," in *Proc. ICLR*, 2021.
- [24] T. Akiba, S. Sano, T. Yanase, T. Ohta, and M. Koyama, "Optuna: A next-generation hyperparameter optimization framework," in *Proc. ACM SIGKDD*, 2019, pp. 2623–2631.
- [25] B. K. Iwana and S. Uchida, "An empirical survey of data augmentation for time series classification with neural networks," *PLOS ONE*, vol. 16, no. 7, e0254841, 2021.
- [26] H. Zhang, M. Cissé, Y. N. Dauphin, and D. Lopez-Paz, "mixup: Beyond empirical risk minimization," in *Proc. ICLR*, 2018.
- [27] J. S. Richman and J. R. Moorman, "Physiological time-series analysis using approximate entropy and sample entropy," *Am. J. Physiol. Heart Circ. Physiol.*, vol. 278, no. 6, pp. H2039–H2049, 2000.
- [28] C. Bandt and B. Pompe, "Permutation entropy: A natural complexity measure for time series," *Phys. Rev. Lett.*, vol. 88, no. 17, art. 174102, 2002.
- [29] B. Hjorth, "EEG analysis based on time domain properties," *Electroenceph. Clin. Neurophysiol.*, vol. 29, no. 3, pp. 306–310, 1970.
- [30] T. Higuchi, "Approach to an irregular time series on the basis of the fractal theory," *Physica D*, vol. 31, no. 2, pp. 277–283, 1988.
- [31] M. Rostaghi and H. Azami, "Dispersion entropy: A measure for time-series analysis," *IEEE Signal Process. Lett.*, vol. 23, no. 5, pp. 610–614, 2016.
- [32] Task Force of the ESC and NASPE, "Heart rate variability: Standards of measurement, physiological interpretation, and clinical use," *Circulation*, vol. 93, no. 5, pp. 1043–1065, 1996.
- [33] M. Brennan, M. Palaniswami, and P. Kamen, "Do existing measures of Poincaré plot geometry reflect nonlinear features of heart rate variability?," *IEEE Trans. Biomed. Eng.*, vol. 48, no. 11, pp. 1342–1347, 2001.
- [34] E. Perez, F. Strub, H. de Vries, V. Dumoulin, and A. Courville, "FiLM: Visual reasoning with a general conditioning layer," in *Proc. AAAI*, 2018.
- [35] M. Ibarra-Manzano *et al.*, "ITSC: Inter-turn short-circuit fault dataset for three-phase induction motors," GitHub repository, <https://github.com/ibarram/ITSC>.
- [36] R. Cunha, "MIT short-circuit flux and current signals dataset," Kaggle, 2021, <https://www.kaggle.com/datasets/rebecacunha/mit-short-circuit-flux-and-current-signals>.
- [37] A. Vaswani *et al.*, "Attention is all you need," in *Proc. NeurIPS*, 2017.

- [38] S. Bai, J. Z. Kolter, and V. Koltun, “An empirical evaluation of generic convolutional and recurrent networks for sequence modeling,” *arXiv:1803.01271*, 2018.
- [39] A. Sharma *et al.*, “Recurrence-plot-based severity assessment of stator inter-turn faults using ensemble learning,” *Measurement*, 2023.